

# 8. Tutorial on inverse ray shooting

JORGE JIMÉNEZ-VICENTE

Magnification maps are an essential tool in microlensing studies. Their calculation is based on very simple principles and it is therefore quite straightforward to implement. This tutorial is intended to show how these calculations are done by using a basic ray-shooting procedure. The tutorial assumes some basic knowledge of any programming language, but no previous knowledge of the specific language used here or experience with gravitational lensing computations is needed. The very basics of gravitational lensing are also implicitly assumed at some points. From the computational point of view, the tutorial covers topics ranging from the simplest ray-shooting program for generating images of an object through a simple lens system to the production of magnification maps for quasar microlensing. Source size effects and how to deal with them are also briefly discussed. We finish by also briefly discussing the main improvements that have been introduced into this technique to make calculations faster.

## 8.1. Introductory remarks

It is particularly difficult to put into writing what was intended as a couple of practical sessions on inverse ray-shooting techniques during the Winter School. Unfortunately, there is not much choice but to illustrate the ideas with some code. At this point, a choice has to be made whether to use pseudocode or to choose a given programming language. I have chosen the latter option here in the hope that the reader may actually use the code snippets presented here straight away and be able to produce some useful programs from it. During the lectures, the **Python** programming language was used for the tutorial. I chose it for a couple of reasons that are enumerated below. I have therefore also used **Python** here for presenting the techniques introduced in this tutorial. Nevertheless, the reader is not expected to know any **Python** in advance as I shall introduce all the required information on **Python** syntax in Section 8.2. Readers should therefore find no problem in following the explanations and/or in translating the code to their favourite programming language. Finally, these lectures were intended as practical sessions and, as such, the focus is on producing some useful code and on understanding the key steps in that process. Performance will be dealt with at some point in order to introduce the necessary **Python** ingredients to make the programs fast enough to produce results in a reasonable time. But, while performance is not something I have been particularly worried about in this tutorial, it is indeed important for real calculations. Therefore, although the reader should not expect to find here a thorough description of the many advanced techniques developed for improving performance in these kinds of calculations, some of the most successful and widely used techniques are briefly discussed at the very end of this tutorial.

## 8.2. A short introduction to Python, Numpy, and matplotlib

As mentioned above, **Python** will be used throughout this tutorial as the programming language to illustrate inverse ray-shooting techniques in gravitational lensing. The main reasons for choosing **Python** for this purpose are:

*Astrophysical Applications of Gravitational Lensing*, ed. E. Mediavilla et al. Published by Cambridge University Press. © Cambridge University Press 2016.

- (i) **It is free and easily available for most operative systems.** **Python** is available in almost any operative system running on almost any hardware, from supercomputers to mobile phones and gaming consoles. In particular, it is available in Linux/Unix in most (all?) distributions/flavours, Windows, and Mac OS/OS X. This is certainly a very important aspect. Anybody can get **Python** and install it on virtually any computer for free in just a few minutes. This is indeed most convenient for a tutorial like this.
- (ii) **It is an interpreted language.** This is also a very handy property of **Python** (or more precisely of its standard implementation<sup>†</sup>). Therefore, after making any change to the code, it can be executed immediately without the need to recompile and with no more ado.
- (iii) **It is easy to read/write and therefore easy to learn.** Indeed, in many aspects there is not much difference between pseudocode and **Python** code. Being a dynamic typed language, the interpreter takes care of data types at run time, and the user need not worry about type declaration but is able instead to concentrate on other things. Memory management is also something the user may forget about for most purposes. It is therefore very well suited for teaching purposes.
- (iv) **It is versatile and powerful (especially when combined with existing libraries).** The **Python** standard library is by itself a powerful tool with endless possibilities,<sup>‡</sup> but there are also now plenty of external libraries which make a programmer's life even easier. Being a general-purpose programming language, **Python** can access plenty of libraries for many specific purposes. Here, we make use of four of them, since they are very handy in scientific computing: **Numpy** operates in an efficient way with arrays, **Scipy** for some scientific calculations and/or image processing, **matplotlib** for plotting and images, and **pyfits** for dealing with fits files.

If I were interested in convincing you to use **Python**, I could give you a few more reasons, such as it is really fun to use, or that it is a good investment in the mid-term for a scientist in general and for an astronomer in particular. I could even try to impress you by telling you that **Python** is behind sites as well known as YouTube, and that it is widely used in scientific institutions as prestigious as NASA and CERN. But that is not the case, so I shall stop the marketing here. You will find plenty of information on the official webpage (<http://www.python.org>).<sup>§</sup>

**Python** can be used for procedural/functional programming, as well as for object-oriented programming. Although I may occasionally use some built-in object-oriented methods, I shall follow a fully procedural/functional paradigm here. The code presented here will probably not be very *Pythonic* and will violate most of the guiding principles in *The Zen of Python*. I hope the *Benevolent Dictator For Life* will forgive me for this. Not so sure about his acolytes.

<sup>†</sup> In fact, **Python** is *not* interpreted. In the standard implementation, source code is indeed compiled to bytecode, which is executed in a virtual machine. From the user's perspective these processes are completely hidden so that, in fact, it feels as though it is indeed an interpreted language.

<sup>‡</sup> **Python** enthusiasts refer to this fact with the phrase 'batteries included'.

<sup>§</sup> The code in this tutorial has been written to work with **Python 2.6** or **2.7**. In 2008 **Python 3.0** was released: this version is not backwards compatible with the 2.x versions. Although versions of all the necessary libraries for this tutorial are available for **Python 3.1** and **3.2**, I have no experience with it and cannot recommend or discourage its use.

### 8.2.1 Python in a (pine) nut shell

**Python** can be used in essentially two different ways: interactively or by using scripts. The **Python** interpreter can be invoked by typing **python** at the command line prompt. Once started, we can begin to type commands. **IPython** is a fancy shell for interactive **Python** which is very convenient.

Alternatively, a script can be created with an editor. Usually, **Python** programs are indicated by a **.py** extension. A script named **script.py** can be executed as an argument of the **python** command with **python script.py**. As **Python** scripts can themselves have arguments, these should follow the script name. There are some useful development environments for **Python**, **IDLE** being widely used.

We very quickly pass over the most basic rudiments of **Python** here to be able to follow the rest of the tutorial. The reader interested in more detail is referred to *The Python Tutorial*,<sup>†</sup> or to the one written by Shipman (2011). Both are excellent places to look for more detailed information.

#### 8.2.1.1 Variables and data types

We have already stated that **Python** is dynamically typed, meaning that variable names do not have an associated data type. Inside **Python**, objects do indeed have a well-defined type, but we can refer to them by any name at any time. Types are linked to the objects, not to their names. Therefore, a variable name **a** can be used to label an integer and, within the same session or program, be reused to label a string, float, or any other object. This is extremely handy, as we do not usually have to care much about variable declaration, but this also entails certain risks. If we mistype a variable name when trying to operate on some object, **Python** will not complain but will do whatever we asked it to do if that is possible. This type of error is sometimes not easy to debug, so one has to be particularly careful.

Creating variables and/or assigning them values in **Python** is therefore straightforward<sup>‡</sup>:

```
a = 1 # a is an integer
a = 1.3 # a is now a float
a = 'Hallo!' # a is a string
print a[3] # We can refer to a character within a string
#But we cannot change it -> a[3]='k' is forbidden !!. Strings are not mutable.
#In Python indices start from 0 for the first element. C-like
```

There are also three different groups of objects: lists, dictionaries and tuples:

```
a = [1,3,4,'Hallo'] # a is a list. Elements in a list do not need to be of the
    same type
a[0] = 7 # This is the first element. Lists are mutable
a[-1] = 'Bye' # Negative indices start from the end. This is the last element
print a[1:-1] # This is [3,4]. This is called a slice
a = [[1,2],[3,4]] # Elements of list can also be lists. But a list of lists is
    not an array!!
a = {'mass':3.2,'vel':3.2e5,3:4} # Dictionary. Unordered named list.
a['mass'] = 6.7 # Elements of dictionaries are referred by keyword, not by
    position. Dictionaries can be changed. They are mutable
a = (1,3,5,'siesta') # A tuple. Similar to a list but non mutable
# a[0]=5 This is not permitted if a is a tuple.
```

<sup>†</sup> <http://docs.python.org/2/tutorial/index.html>.

<sup>‡</sup> The **#** character is used in **Python** to comment the rest of the line.

### 8.2.1.2 Flow control, functions and input/output

Probably the most distinctive hallmark of **Python** is the way in which code blocks are indicated. In **Python** code blocks do not use braces, parentheses or brackets,<sup>†</sup> but are indicated by the level of indentation. This means that in **Python**, indentation is much more than just a nice way to show code structure: it is the code structure itself. Although it may seem strange at first, one gets used to it quite quickly. Although not mandatory, it is customary to use four spaces per indentation level. In general, it is not safe to mix tabs and spaces, so try *always* to stick to four spaces per indentation level to avoid problems.

Among the different structures to organize program flow, we present here only three: **for** and **while** loops and **if ... else** blocks. A **for** loop looks something like:

```
for i in range(2,5):                # variable i runs from 2 to 5 (excluding 5)
    print i, "bottles of wine are too many"
```

Explicit loops are not particularly fast in **Python**, and this becomes more pronounced when loops are nested. This is why more efficient alternatives such as comprehension lists or libraries such as **Numpy** to deal with arrays are preferred performance-wise. Loops can also be made with a **while** and a condition check:

```
i=0
while( i < 10 ):
    print i
    i += 1
```

We can check for conditions with **if ... elif ... else** structures:

```
if (a[i] == 7): # Next block is executed if condition is fulfilled
    print a[i], "is equal to seven"
    a[i] += 1
elif (a[i] == 5): # If this condition is fulfilled then
    print a[i], "is five"
    a[i] -= 1
else: # If none of the above is fulfilled
    print a[i], "is not 5 or 7"
```

There are some more control flow statements, but we do not use them here. These flow control structures can be combined at will to produce the desired outcome.

**Python** has around 80 built-in functions that can be used straight away, such as **abs()**, **help()**, **raw\_input()**, etc. But the standard library that comes with **Python** has many more functions, and endless libraries are available in **Python**. Functions are usually organized into modules that contain several related functions. For namespace sanity, modules and submodules are often organized into packages. Modules can be imported into a **Python** program with the **import** statement, and functions within that module can then be used in that program. For example, we can import and use functions in the **random** module from the standard library with:

```
import random # Here the module is imported
print random.random() # Print a float in [0.0,1.0). Look at the syntax module.func
```

The reason to refer to the function in this way, with the module name in front, is to avoid namespace mixing. We could have another **random()** function in another module and we

<sup>†</sup> Try `from __future__ import braces` within **Python** and see what happens.

could use both while still preventing confusion. We can use an alias for the module with the **as** clause:

```
import random as r # Import the random module with alias r
print r.random()
```

Indeed, we could completely avoid writing the module name if we import the module with:

```
from random import * # Names in random imported to current namespace. BEWARE!!
print random()
```

But this is indeed a dangerous practice due to namespace mixing and potential confusion and is therefore not advisable. We could even import just one or a few names/functions into the current namespace with:

```
from random import random # Import random() to current namespace. Somewhat safer
    but still BEWARE!!
print random()
```

Of course, we can do much more than use already made functions: we can create our own functions/modules. A function is created with the **def** clause:

```
def factorial( n ): # def function ( arg1, arg2, arg3, ... )
    if n < 1: # Base case. Notice the indentation of the whole function code
        return 1
    else:
        return n * factorial( n - 1 ) # Python functions can be recursive !!
```

Several functions can be defined on the same file as a module that can be imported like any other module.

A **Python** script can also have command line arguments. These can be accessed via the variable **argv** of module **sys**. We may have something like:

```
import sys
nargs=len(sys.argv)-1 # Number of arguments. argv[0] is the script name.
for i in sys.argv[1:]: # Loop over arguments (excluding script name)
    print i
```

Of course, there are more complex ways to parse arguments, but this should be enough for our needs here.

Files in **Python** are objects that can be created with the **open()** function. The first argument of **open()** is the file name. A second argument in **open()** can be used to set whether the file is open for reading ('r'), writing ('w'), appending ('a'), etc. If not set, the file is open for reading. We can read a line with **readline()**. Indeed, files are iterable objects, so we can read a file with:

```
for i in open(filename): # A file is iterable. Loop over lines.
    print i              # i is a string containing a line in file, not a number.
```

Or we can write to it with:

```
f=open(filename,'w') # BEWARE, opening in 'w' mode will erase the file contents.
f.write(string)      # Write string to file
f.write(str(float))  # Write a float. Convert to string first with str() !!
```

Again, there are more sophisticated ways to deal with input/output, but this is the minimum required.

### 8.2.2 Numpy and matplotlib

We have already said that the standard **Python** implementation is not particularly fast, and this weakness becomes patent in scientific computing when calculations with large arrays and/or nested loops are needed. Fortunately, there is a library called **Numpy** that eases these deficiencies and makes array manipulation very natural and fast. **Numpy** can be imported in the usual way, but in order to avoid namespace mixing, we will invariably import it with **import numpy as np**. Arrays are **Numpy** objects that can be created with **Numpy** functions from scratch or from pre-existing **Python** objects. Let us have a look at an example to see how natural it is to work with arrays within **Numpy**:

```
import numpy as np          # Import numpy module
z=np.zeros(5)              # z is an array of 5 elements filled with zeros
o=np.ones((3,3))          # o is a 3x3 array filled with ones
l=np.linspace(3,5,21)      # Array with 21 elements from 3 to 5 (included)
a=[1,2,3,4,5]             # a is a python list
a=np.array(a)              # Convert a list into an array
c=a+b                      # c,a and b are arrays.
c=3*a                     # Multiply every element by 3
c=a*b                     # Per element multiplication
c=np.exp(a)                # Apply function to every element
b=np.ones(3)               # b is broadcasted to right size before addition
c=o+b                      # b is broadcasted to right size before addition
y,x=mgrid[0:3,0:3]        # x and y are 3x3 arrays with column and row values
r=np.sqrt((x-1)**2+(y-1)**2) # Array with distance to pixel 1,1
```

In **Numpy**, arrays are stored in row order as in **C** (unlike **FORTRAN** or **IDL**, which store arrays in column order). This means that, as we move linearly through memory in an array, the rightmost index changes the fastest. This is an important fact to be able to use the memory cache efficiently in many cases. We must remember that arrays do multiply on a per basis element, not like matrices and/or vectors. If we need that behaviour, we should use **matrix** objects. Arrays can be converted into matrices with the command **np.matrix(arr)**, where **arr** is assumed to be an array. We do not use matrices in this tutorial.

To end this section, we briefly mention how to make simple plots in **Python**. This is indeed very easy thanks to the **matplotlib** library. Let us illustrate it with a simple example:

```
import numpy as np
import matplotlib.pyplot as plt    # Import the pyplot module from matplotlib
x=np.linspace(-5,5,101)           # Array of 101 elements from -5 to 5
y=np.exp(-x**2/2.)                # Gaussian function
plt.plot(x,y)                     # Plot x vs y
plt.plot(x,y,'+g')                # Same using green + symbols
plt.show()                        # Show the plot
```

The plot is shown in an interactive window that allows interaction, including zooming in and out, and even saving the plot into a file in different formats (gif, eps, pdf, etc.). The plot function in **matplotlib** has plenty of options that we will not cover here. The interested reader can have a look at the official **matplotlib** documentation

at <http://matplotlib.org/contents.html>. We can also make subplots in a page with the **subplot** command. The line **plt.subplot(nrows,ncols,nplot)** creates **nrows**×**ncols** subplots, **nplot** being the active plotting axes. If **nrows**, **ncols** and **nplot** are all smaller than 10, we can skip the commas. Thus **plt.subplot(221)** sets the upper left corner as the active plotting axes in a group of four subplots set in two rows and two columns. We can also plot images with the **plt.imshow(array, vmin=min, vmax=max)** command. We remind the reader to use the **plt.show()** command to show all the plots unless the interactive mode is being used.

### 8.3. Inverse ray-shooting basics

The gravitational lens equation can be expressed in dimensionless units as (Schneider, Ehlers and Falco 1999)

$$\mathbf{y} = \mathbf{x} - \boldsymbol{\alpha}(\mathbf{x}), \quad (8.1)$$

where  $\mathbf{y}$  and  $\mathbf{x}$  are the (2D vector) coordinates at the source and lens plane respectively, and  $\boldsymbol{\alpha}(\mathbf{x})$  is (also a 2D vector) called the scaled, or reduced, deflection angle. Solving this equation means being able to invert it and to obtain the position of the image(s)  $\mathbf{x}_i$  corresponding to a source point  $\mathbf{y}$ . This equation can be interpreted as describing the trajectory of a ray of light going *backwards* from location  $\mathbf{x}$  at the lens plane to location  $\mathbf{y}$  at the source plane. For this reason it is sometimes referred to as the ray-trace equation of gravitational lensing. As light rays are invariant under time reversal, tracing them forward (as real photons are travelling) or backwards makes no difference from the mathematical point of view. We assume here the *thin lens* approximation, in which the deflection of light takes place mostly in a very narrow region of space which is very small compared to the distances between lens and source ( $D_{LS}$ ), lens and observer ( $D_L$ ) and, of course, source and observer ( $D_S$ ).

This equation is, in general, highly non-linear due to the dependence of  $\boldsymbol{\alpha}$  on the position on the lens plane. Indeed, it can only be solved analytically in a few simple cases. As we know already, in some circumstances, this equation can even produce multiple images of a source (so that there are several values of  $\mathbf{x}$  that are solutions to equation (8.1) for a given value of  $\mathbf{y}$ ). When the lens and/or source are complex, equation (8.1) can only be inverted by numerical methods. Here, our ultimate goal is to be able to calculate magnification maps for quasar microlensing, but in the meantime, we will also apply the procedure to simpler cases.

As we have indicated above, the interpretation of the lens equation as a ray-trace equation is very natural. Therefore, the technique of shooting rays from the observer plane backwards, deflecting them at the lens plane, and tracing them to the source plane is a very simple but effective way to solve the lens equation for a given region of the image plane. This technique is called inverse (or backwards) ray shooting (IRS) and was introduced by Schneider and Weiss (1986, 1987) and Kayser, Refsdal and Stabell (1986) to calculate magnification maps. Here, we introduce this technique from the most simple case of a point lens (even the trivial case of no lens will be explored to set up the main structure of the code) to the complex case of quasar microlensing.

The underlying idea is, therefore, to study a certain region of the image/observer plane by tracing backwards the trajectories of a set of light rays and looking what is in the source plane at the position of the ray hit. The trajectory between observer and lens plane is a trivial straight line, as there is no deflection. At each point in the lens plane, the rays are deflected an amount given by the scaled deflection angle. From the lens plane to the source plane, the trajectory is again a straight line, and therefore there is



no further change in the coordinates. The convergence/divergence of light rays between lens and source planes is already implicitly taken into account by the different angular scales used at the lens and source plane in equation (8.1).

The algorithm goes as follows:

- (i) Divide the source and image plane regions under study into cells/pixels. Let  $2x_l$  and  $2y_l$  be the sizes of the regions at the lens and source planes that we are going to consider. We divide both planes into  $n_x$  and  $n_y$  cells (or pixels) respectively. Therefore, the size of a cell is  $x_s = 2x_l/(n_x - 1)$  and  $y_s = 2y_l/(n_y - 1)$  in the lens and source plane respectively.<sup>†</sup> The input parameters are therefore the size and resolution of the regions at the source and image plane.
- (ii) Set the surface brightness at the source plane according to the source model. We should know the surface brightness at the source plane. To start with something easy, let us define a Gaussian circular source of sigma *rad* located at pixel  $(x_1, y_1)$ . Let us first write a function **gcirc(ny,rad,x1,y1)** that returns an array of  $n_y \times n_y$  with such a source. We can create a module **source.py** like this:

```
import numpy as np
def gcirc(ny,rad,x1=0.0,y1=0.0):          # Default position is 0,0
    x,y=np.mgrid[0:ny,0:ny]
    r2=(x-x1-ny/2)**2+(y-y1-ny/2)**2
    a=np.exp(-r2*0.5/rad**2)
    return a/a.sum()
```

The input parameters here are the location, size and structure of the source.

- (iii) Trace rays back by deflecting them at the lens plane according to the lens equation (8.1). Now, we just need to trace back rays from the centre of each of our cells at the lens plane and see where these rays hit at the source plane.
- (iv) Find the cell/pixel at the source plane where the ray hits. If the ray hits at a spot in the source plane where there is a bright source, then the value of that cell in the image/lens plane is set equal to the surface brightness of the source at that position. Even if there is a single source, the bending of light rays by the effect of the lens may end up producing several images of it. This would happen if rays starting from different positions at the image plane end up hitting the same spot of the source plane.

The last two steps in this algorithm involve conversion from pixels to coordinates (at the lens plane) and back to pixels (at the source plane after deflection). To start with the easiest case, let us just explore the case with no lens and therefore no deflection of the light rays. The reader may think that there is no interest in such a case, but this way we focus only on the general set-up of the program, independently of the lensing calculations. The input parameters are the location and structure (mass, shape, model) of the lens (which, in this first example, will be set to a no-lens model).

An example of code for doing this is given in Table 8.1.

It is an interesting exercise to play around a bit with the values of the input parameters, in particular the size of the source and of the source and image regions and the resolution in these planes to see what the effect is. For example, it does not make much sense having a resolution at the image plane much higher than at the source plane. We can also see

<sup>†</sup> We are implicitly assuming here that  $\pm x_l$  and  $\pm y_l$  are indeed the coordinates of the centre of the extreme cells and not the coordinates of the cell's edges. The real size of the regions in the lens and source plane are indeed  $2x_l + x_s$  and  $2y_l + y_s$ . The reason to do it like this is that when pixels are numbered with  $i$  from 0 to  $n_x - 1$ , coordinates of cell centres defined as  $x_1 = -x_l + i * x_s$  run from  $-x_l$  to  $x_l$ . Same for coordinates in the source plane.



TABLE 8.1. Code for `irs0.py`

---

```

import numpy as np
import matplotlib.pyplot as plt
import source as s

nx=401                                     # Number of pixels in image plane
ny=401                                     # Number of pixels in source plane
xl=2.                                     # Half size of image plane covered (in "Einstein" radii)
yl=2.                                     # Half size of source plane covered (in Einstein radii)

xs=2.*xl/(nl-1)                           # pixel size on the image map
ys=2.*yl/(ny-1)                           # pixel size on the source map

# Source parameters
xpos=0.0                                  # Source position. X coordinate
ypos=1.0                                  # Source position. Y coordinate
rad=0.1                                  # Radius of source
ipos=int(round(xpos/ys))                   # Convert source parameters to pixels
jpos=int(round(-ypos/ys))
rpix=int(round(rad/ys))
a=s.gcirc(ny.rpix,jpos,ipos)              # This is a circular gaussian source
b=np.zeros((nx,nx))                      # This is the image plane

# This is the main loop over pixels at the image plane
for j1 in range(nx):
    for j2 in range(nx):
        x1=-xl+j2*xs # Convert pix to coords on image
        x2=-xl+j1*xs
        y1=x1-0.0    # Deflect X coordinate
        y2=x2-0.0    # Deflect Y coordinate
        i2=int(round((y1+y1)/ys)) # Convert coordinates to pixels
        i1=int(round((y2+y1)/ys))
        # If deflected ray hits a pixel within source then set image
        # to brightness on that pixel
        if ((i1 >= 0) and (i1 < ny) and (i2 >= 0) and (i2 < ny)):
            b[j1,j2]=a[i1,i2]

# Plot stuff
plt.subplot(121)
plt.imshow(a,extent=(-yl,yl,-yl,yl))
plt.subplot(122)
plt.imshow(b,extent=(-xl,xl,-xl,xl))
plt.show()

```

---

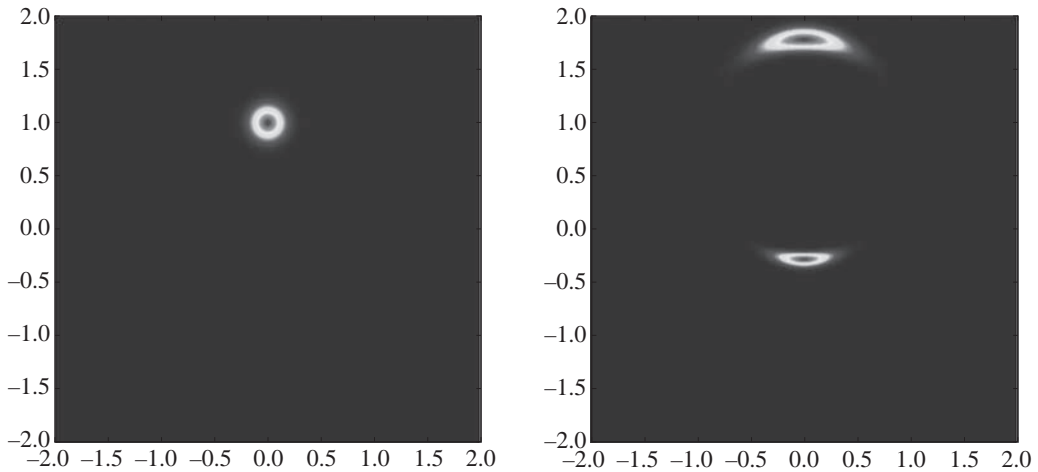


FIGURE 8.1. Source (left) and image (right) planes. Point lens of mass = 1 located at (0,0.5). Source is a circular Gaussian source at  $x = 0$ ,  $y = 1$ .

that if the source is small and the resolution at the image plane is small, we may miss the image or have a very unrealistic one. Thus, source and image plane should both be conveniently sampled.

### 8.3.1 *Image(s) for a simple lens: the point source lens*

The next step in our path towards producing magnification maps is the study of the case of a point source lens of mass  $m_l$  located at  $\mathbf{x}_d = (x_d, y_d)$ . In this case, we set the deflection of the coordinates as  $\alpha = m_l(\mathbf{x} - \mathbf{x}_d)/|\mathbf{x} - \mathbf{x}_d|^2$  where  $\mathbf{x} = (x_1, x_2)$  are the coordinates of the ray at the image plane. As we did with the source, we can create a function that calculates this deflection for such a ray in a module called **lens.py**. The function will look something like<sup>†</sup>:

```
def Point(x1,x2,x1l,x2l,ml):      # Point lens of mass ml at x1l,x2l
    x1ml=(x1-x1l)                # Distance along x axis of ray to lens position
    x2ml=(x2-x2l)                # Distance along y axis of ray to lens position
    d=x1ml*x1ml+x2ml*x2ml+1.0e-12 # Distance between ray and lens squared
    y1=x1-ml*(x1-x1l)/d          # Lens equation for x coordinate
    y2=x2-ml*(x2-x2l)/d          # Lens equation for y coordinate
    return (y1,y2)
```

We see that the function returns a tuple  $(y_1, y_2)$  with the coordinates of the deflected ray at the source plane. After importing the module with **import lens as l** at the beginning of our previous program, we can then set the lens parameters  $x_d, y_d, m_l$ , and we can then deflect rays by calling this function as **y1,y2=l.Point(x1,x2,xd,yd,ml)**. If we locate our source at point (0,1) and the lens of mass  $m_l = 1$  at position (0,0.5), the result is shown in Figure 8.1. We can see that two images appear with a separation  $d \approx 2$ . Indeed, that is exactly the theoretical result. If source and lens are aligned, the image is a nice Einstein ring of radius  $\theta_E = 1$ . The reader is encouraged to play around a bit with the input parameters, in particular with the relative positions of lens and source. The magnification

<sup>†</sup> The  $10^{-12}$  added to the distance squared is a very inelegant yet simple way to avoid division by zero when lens and source are aligned. Sorry about that!

can be calculated as the ratio of the flux in both planes (taking into account the pixel size on each plane).<sup>†</sup>

## 8.4. Playing around with lenses and sources

Once we are familiar with the simplest case, it is time to explore other popular lens models. We investigate some well-known models which are described in Chapter 8 of Schneider et al. (1999). The procedure is very similar to the one we have followed with the point source lens by creating a function that calculates the deflection of a given ray at location  $(x_1, x_2)$  in the lens plane. As an example, we explore three well-known cases:

- (i) **The binary lens** is a straightforward extension of the previous case by just adding a new lens of mass  $m_{2l}$  at position  $(x_{2d}, y_{2d})$ . The deflection angle for a ray at location  $\mathbf{x} = (x_1, x_2)$  is given by

$$\boldsymbol{\alpha} = m_{1l} \frac{(\mathbf{x} - \mathbf{x}_{1d})}{|\mathbf{x} - \mathbf{x}_{1d}|^2} + m_{2l} \frac{(\mathbf{x} - \mathbf{x}_{2d})}{|\mathbf{x} - \mathbf{x}_{2d}|^2}. \quad (8.2)$$

The binary lens is a very well studied case (Schneider and Weiss 1987) and is widely applied, for example, in modelling microlensing events. The two relevant parameters are the mass ratio of the lenses and their separation.

- (ii) **The Chang–Refsdal lens** is a simple extension of the point lens that includes the effect of a smooth background gravitational field by including a quadrupolar term in the deflection angle caused by the shear due to the background gravitational field of a nearby source. It was introduced by Chang and Refsdal (1979, 1984) to address the lensing of a star in the background potential of a galaxy. The effect of the background potential is included via its convergence  $\kappa$  and shear  $\gamma$  so that the deflection angle (in the coordinate system in which one of the axes runs along the direction of the shear) is

$$\boldsymbol{\alpha} = \begin{pmatrix} \kappa + \gamma & 0 \\ 0 & \kappa - \gamma \end{pmatrix} \mathbf{x} + m_l \frac{(\mathbf{x} - \mathbf{x}_d)}{|\mathbf{x} - \mathbf{x}_d|^2}. \quad (8.3)$$

- (iii) **The singular isothermal sphere (SIS)** is characterized by its Einstein radius,  $\theta_E = 4\pi (\sigma_v^2/c^2) (D_{LS}/D_S)$ , where  $\sigma_v$  is the one-dimensional velocity dispersion. The deflection angle for this lens system is given by

$$\boldsymbol{\alpha} = \theta_E \frac{(\mathbf{x} - \mathbf{x}_d)}{|\mathbf{x} - \mathbf{x}_d|}. \quad (8.4)$$

It is a simple model that is often used to mimic the lensing properties of galaxies and/or clusters of galaxies.

The singularity at the origin can be removed by introducing a finite core that changes the lensing properties very close to the lens position. That model is called the non-singular isothermal sphere.

Including these lenses in our **lens.py** module is quite straightforward following the example of the point source lens. The module would look something like Table 8.2.

Any other lens models (such as more than two point lenses, singular and non-singular isothermal spheres, or ellipsoids with or without external shear, etc.) can be included in a similar way. We can now test our former IRS code by playing with these new lenses.

<sup>†</sup> Recall that this is a magnification *averaged* over the whole source. This should not be confused with the value of the magnification of a point source lens at the position of the source.

TABLE 8.2. Code for lens.py

---

```

import numpy as np

def Point(x1,x2,x1l,x2l,ml):
    x1ml=(x1-x1l)
    x2ml=(x2-x2l)
    d=x1ml*x1ml+x2ml*x2ml+1.0e-12          # Add a tiny number to avoid division by zero
    y1=x1-ml*(x1-x1l)/d                  # Lens equation for x coordinate
    y2=x2-ml*(x2-x2l)/d
    return (y1,y2)

def TwoPoints(x1,x2,x1l1,x2l1,x1l2,x2l2,ml1,ml2):
    x1ml1=(x1-x1l1)
    x2ml1=(x2-x2l1)
    x1ml2=(x1-x1l2)
    x2ml2=(x2-x2l2)
    d1=x1ml2*x1ml1+x2ml1*x2ml1+1.0e-12    # Add a tiny number to avoid division by zero
    d2=x1ml2*x1ml2+x2ml2*x2ml2+1.0e-12
    y1=x1-ml1*(x1-x1l1)/d1-ml2*(x1-x1l2)/d2 # Lens equation for x coordinate
    y2=x2-ml1*(x2-x2l1)/d1-ml2*(x2-x2l2)/d2
    return (y1,y2)

def ChangRefsdal(x1,x2,x1l,x2l,ml,k,g):
    x1ml=(x1-x1l)
    x2ml=(x2-x2l)
    d=x1ml*x1ml+x2ml*x2ml+1.0e-12          # Add a tiny number to avoid division by zero
    y1=x1*(1.0-k-g)-ml*(x1-x1l)/d
    y2=x2*(1.0-k+g)-ml*(x2-x2l)/d
    return (y1,y2)

def SIS(x1,x2,x1l,x2l,k):
    x1ml=(x1-x1l)
    x2ml=(x2-x2l)
    d=np.sqrt(x1ml*x1ml+x2ml*x2ml+1.0e12)  # Add a tiny number to void division by zero
    y1=x1-k*(x1-x1l)/d                    # Lens equation for x coordinate
    y2=x2-k*(x2-x2l)/d
    return (y1,y2)

```

---

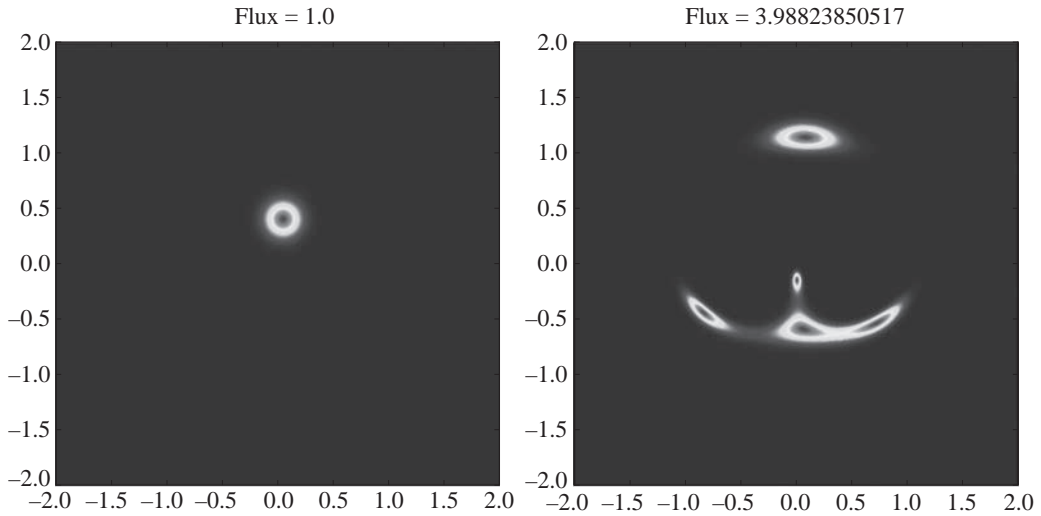


FIGURE 8.2. Source (left) and image (right) planes. Binary point lens of mass 0.5 located at  $(0.5, 0.0)$  and  $(-0.5, 0.0)$ . Source is a circular Gaussian source at  $x = 0.05$ ,  $y = 0.4$ .

For example, we can try to reproduce the configurations shown in Schneider and Weiss (1986). Figure 8.2 shows a situation similar to that of Figure 6b in Schneider and Weiss (1986).

The reader is encouraged to try these lens models or any other. Change the input parameters for the lens and/or source and try to reproduce the known theoretical results about number of images, parity, etc. It is also instructive to use different sources, such as the image of a real galaxy, or a field of galaxies. This can be read from a fits file using the `getdata` function from the `pyfits` library of **Python**. We can include the following function in our `source.py` module to read a fits image as source:

```
from pyfits import getdata
def fitsim(filename):
    a=getdata(filename)           # Read the file
    if (len(a.shape) > 2): a=a[0]  # Take first plane if there are many.
    return (1.0*a)/a.sum()        # Return the normalized image
```

The source can then be read from our IRS code with:

```
a=s.fitsim('Edgeon2.fits')  # Read file 'Edgeon2.fits'
ny=a[0].size                # Number of pix in souce plane set to image size
```

An example of the result of our code for an edge-on galaxy lensed by a SIS with some shear is shown in Figure 8.3.

This **Python** code is useful and may allow us to gain insight into some lens systems, but we have to accept that it is not very efficient. There are two nested loops (over rows and columns of the image plane) and, if there are many lenses, there will be a third one over lenses. **Python** is, in general, not very efficient in running loops, and consequently nested loops are even worse. As the image plane becomes larger and/or we include more lenses, this code becomes very inefficient and, consequently, slow. Fortunately, improving its performance is both easy and elegant by using **numpy** arrays. Instead of running over individual pixels, we can deflect the whole bunch of rays at once by operating on an entire array of coordinates. This way, the nested loop will be dealt with internally by **numpy**,

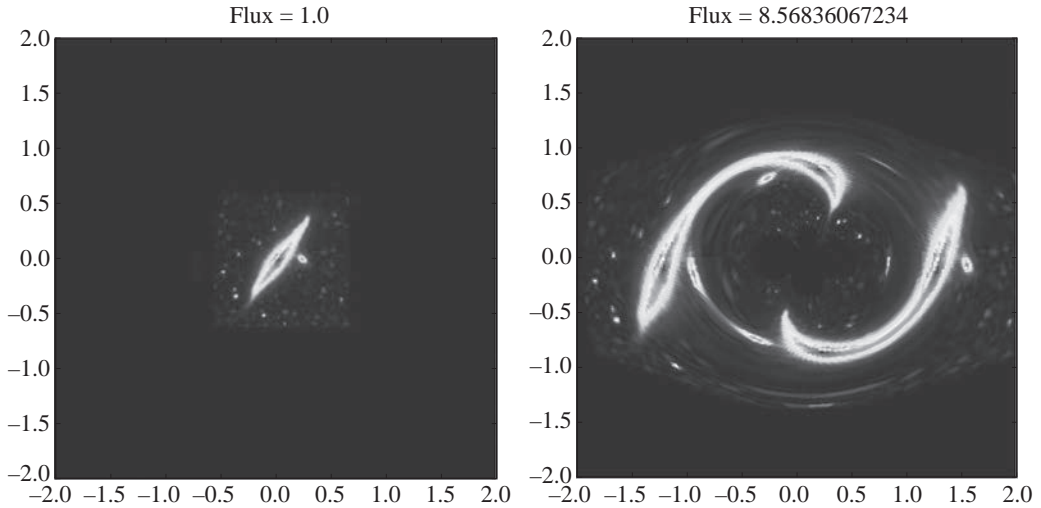


FIGURE 8.3. Source (left) and image (right) planes. SIS lens of  $\theta_E = 1$  and  $\gamma = 0.2$  located at  $(0.1, 0)$ .

which is much more efficient (as **numpy** is already compiled to machine code). And the best news is that, as we were careful enough in writing our functions for the deflection of the different lens systems to use **numpy** functions when necessary, these functions, without any modification, are still as useful when called with full arrays of coordinates as they were with a single ray. We then just need to do some manipulation to create an array with the column and row numbers of the pixels, and remove the loop over pixels by operating on full **numpy** arrays. The code will now look something as shown in Table 8.3.

### 8.5. Magnification maps and light curves

Up to now we have been able to produce images of many different types of sources that are gravitationally lensed by a broad group of lens models. Even if this is instructive and may have some usefulness, we would like to go further and be able to produce magnification maps for our lens systems. A magnification map is one that shows the amount of magnification produced by a certain lens system in a region of the source plane. If, as is usually the case, lens(es) and source move with respect to each other at a certain velocity, the source suffers different magnifications at different times, which are observed as variations in the brightness of the source. A curve of the brightness of a source versus time is called a *light curve*, and contains plenty of information on the lens and/or source systems. Magnification maps are therefore needed to be able to interpret these kinds of observations and to compare them with different possible models. They allow us to determine many properties of the lens system (mass, structure) and/or of the source (size, etc.). In order to calculate magnification maps, we make use of the fact that gravitational lensing does not change the surface brightness of the images, and therefore the magnification is just the ratio between the subtended solid angles of image(s) and source, which is given by the inverse of the determinant of the  $A$  matrix (see Schneider et al. 1999, Section 5.2):

$$\mu(\mathbf{x}) = \frac{d\omega}{d\omega^*} = \frac{1}{\det A(\mathbf{x})}, \quad (8.5)$$

TABLE 8.3. Code for irs1.py	
import numpy as np	
import matplotlib.pyplot as plt	
import source as s	
nx=801	# Number of pixels in image plane
ny=401	# Number of pixels in source plane
xl=2.	# Size of image plane covered (in "Einstein" radii)
yl=2.	# Size of source plane covered (in Einstein radii)
# Lens parameters	
xlen=0.0	
ylens=0.0	
mlens=1.0	
xs=2.*xl/(nx-1)	# pixel size on the mage map
ys=2.*yl/(ny-1)	# pixel size on the source map
# Source parameters	
xpos=0.0	
ypos=0.0	
rad=0.10	
ipos=int(round(xpos/ys))	Convert source parameters to pixels
jpos=int(round(-ypos/ys))	
rpix=int(round(rad/ys))	
a=s.gcirc(ny, rpix, jpos, ipos)	# This is the source plane
b=np.zeros((nx,nx))	# This is the image plane
	# In this version, the main loop is implicit
	# We use operations on numpy arrays instead which is faster
	(cont.)



TABLE 8.3. ( <i>cont.</i> )	
j1,j2=np.mgrid(0:nx,0:nx)	
x1=-x1+j2*xs	# Pix to coord on image x
x2=-x1+j1*xs	# Pix to coord on image y
y1,y2=1.SIS(x1,x2, xlens+0.1,ylens,1.2)	# This line calculates the deflection
i2=np.round((y1+y1)/ys)	
i1=np.round((y2+y1)/ys)	# If deflected ray hits a pixel within source then set image
	# to brightness on that pixel
ind= (i1 >= 0) & (i1 < ny) & (i2 >= 0) & (i2 < ny)	# Now this is an array
	# which is True if the ray
	# hits the source plane.
i2n=i2[ind]	
j1in=j1[ind]	
j2in=j2[ind]	
for i in xrange(np.size(i1n)):	# Loop over pixels that hit the source plane
b[j1in[i],j2in[i]]=a[i1n[i],i2n[i]]	
# Plot stuff including Fluxes in both plane	
fig=plt.figure(1)	
ax=plt.subplot(121)	
ax.imshow(a,extent=(-y1,y1,-y1,y1))	
fa=np.sum(a)	# Flux on source plane
ax.set_title('Flux='+str(fa))	# Set title for subplot 1
ax=plt.subplot(122)	
ax.imshow(b,extent=(-x1,x1,-x1,x1))	
fb=np.sum(b)*(xs**2)/(ys**2)	# Flux on image pl. (taking into account pix size)
ax.set_title('Flux='+str(fb))	# Set title for subplot 2
plt.show()	

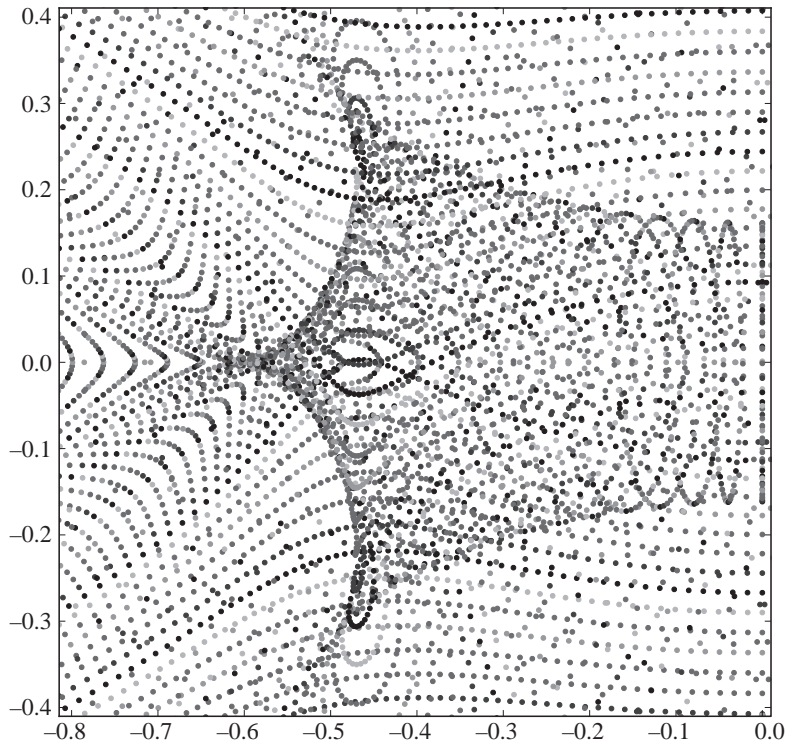


FIGURE 8.4. Ray hits in the source plane for a regular grid of rays at the lens plane. One ray is shot at the centre of every pixel of the lens plane. The lens is a binary lens of equal masses  $m = 0.5$  with a separation of 1.5.

where

$$A(\mathbf{x}) = \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = 1 - \frac{\partial \alpha}{\partial \mathbf{x}}. \quad (8.6)$$

As we are using angular coordinates in the lens and source planes, the solid angle is just an area on those planes. To estimate that ratio of areas we make use of our inverse ray-shooting technique. We shoot a certain (fixed) number of rays per unit area (e.g. per pixel) at the image plane ( $N_r$ ). Now, we collect those rays in cells of the same area at the image plane. If we collect  $N_s(x_1, x_2)$  rays in a cell at coordinates  $\mathbf{x} = (x_1, x_2)$ , then the magnification can be written as

$$\mu(\mathbf{x}) = \frac{dS}{dS_*} = N_s(\mathbf{x})/N_r. \quad (8.7)$$

We must take care with the units in calculating the right size of the shooting and collecting cells to fulfil the second equality of equation (8.7), otherwise we could miss a normalization constant, although this constant can be easily calculated (by checking the unlensed case). Let us see how the method works with a simple example. If we were to shoot rays at the centre of every cell in the image plane and plot the location of the hits at the source plane in the case of a binary lens, we would find something like what is shown in Figure 8.4 (in which rows of rays are marked with the same colour to ease visualization of shape deformation).

We can see that not only do the squares originally formed by contiguous rays get *deformed* into irregular quadrilateral shapes, but very weird things happen in some regions where a large amount of rays bunch together forming funny figures. Those are regions of high magnification according to equation (8.7). Indeed, in those regions of high concentration of rays, the lens transformation is no longer invertible (as those locations correspond to caustics, which are the transformed locations of the critical curves where  $\det A(\mathbf{x}) = 0$ ) and squares are transformed into crossed quadrilaterals. Therefore, a couple of important lessons can be learned from this example:

- Caustics, which are regions of very high (virtually infinite) magnification, have a high concentration of rays and therefore very good statistics for calculating the magnification as described in equation (8.7) via inverse ray shooting. However, critical cells at the lens plane and caustics at the source plane are the places with the most peculiar topological properties of the lens transformation, and particular care must be taken with them for certain purposes.
- Low magnification regions receive few ray hits and consequently have poor statistics in calculating the magnification. The need for a very fine grid of rays to calculate accurate magnification maps comes mainly from the need to have good statistics in these low magnification regions. We will soon see that this is an important fact from which we will be able to profit.

With this experience, we have enough tools to write an IRS code to calculate magnification maps for a simple lens model. The outline of the code could be as follows:

- (i) Prepare a uniform grid of rays at the image plane to be shot backwards towards the source plane. The denser the grid, the better the statistics and the more accurate the final magnification map. Obviously, a compromise must be set between accuracy and computation time. The relevant parameter here (apart from the already mentioned size of the image region and the number/size of pixels) is then the number of rays ( $N_r$ ) shot per (unlensed) pixel at the image plane. Some IRS codes prefer to shoot the rays at random (still keeping an average of  $N_r$  per pixel) instead of on a regular grid. By using a regular grid, we get an accuracy that improves as  $N_r^{3/4}$  (cf. Kayser et al. 1986; Mediavilla et al. 2011), which is better than what would be expected from Poissonian statistics.
- (ii) Loop over the grid of rays and deflect them according to the lens equation. As the lens equation is linear in the deflection angle, if there are several deflectors, loop over the deflectors and add the deflection produced by each of them. Calculate the coordinates of the ray hits at the source plane. It is worth noting here that this step is computationally very demanding, as these loops may be huge if there are many rays to shoot or many deflectors. To get an order of magnitude, let us do some rough calculations. The number of steps in our loops will be

$$N_{\text{steps}} = N_x^2 \times N_r \times N_*, \quad (8.8)$$

where  $N_x$  is the number of pixels on a side of our (for simplicity, square) shooting region (which is determined by the size of our magnification map, as we will see below),  $N_r$  is the number of rays per unlensed pixel, and  $N_*$  is the number of deflectors. A standard magnification map usually needs around 100 rays per pixel, and, if we create a map of  $1000 \times 1000$  pixels, we are shooting  $\sim 10^8$  rays. If we are calculating a microlensing map with a thousand stars as deflectors, we have to calculate  $10^{11}$  deflections! This is the number of steps in our nested loop and is indeed an important figure. No wonder the main efforts aimed at improving the

performance of an IRS algorithm are devoted to reducing this number as much as possible. We will come back to this issue later on.

One last comment is needed here with respect to our **Python** implementation for this nested loop. If we write explicit **Python** loops, which we know are not very efficient, this calculation would take too long. Again, we can speed calculations up by a careful use of **numpy** arrays. We could be tempted again to shoot the whole bunch of rays at once and do the calculations in a single step while keeping only the loop over deflectors. But as we have already seen, the number of rays in a magnification map is usually pretty high ( $\approx 10^8$ ), with a very high risk of running out of memory in most computers if we proceed in this way. We therefore have to lower our expectations a bit and are forced to make a compromise. A pretty good one can be achieved by throwing a whole row (or column) of rays at once.<sup>†</sup> We will use this implementation in our **Python** code below.

- (iii) Set an array of size  $N_y \times N_y$  as our magnification map. Originally, this array's elements are all set to zero. Now, for each ray shoot, we calculate on which pixel at the source plane the ray hits (if at all), and we add 1 to the value at that pixel. After accounting for all ray hits, we divide this array by the number of rays shot at each unlensed pixel in the image plane. The resulting array is the magnification map we were looking for. The size and resolution (pixel size) of this map are determined by the smallest and largest sizes of the sources we are interested in. The pixel size should be smaller than the smallest source (which should span a few pixels). Large maps are usually needed to deal with large sources and/or to have good statistics. Finally, be aware that the size of this map,  $N_y$ , will also determine the size of the shooting region. We will comment briefly on this below.

The resulting code would be something like Table 8.4. The result of running this code is shown in Figure 8.5.

One aspect that must be carefully considered is the choice of the size of shooting region at the image plane. The shooting region must be large enough to contain most (all would be desirable but this is not always possible) rays hitting the magnification map. But it should be small enough to avoid wasting rays that never hit the map. Therefore, how big the shooting region must be depends on the particular lens model. If the shooting region is too small, the map will present *discontinuities* corresponding to regions where there are no ray hits (because there are no rays shot at the appropriate locations in the image plane). The effect of a small shooting region is shown in Figure 8.6 with the same values for  $x_l$  and  $y_l$  as Figure 8.5 but with a Cheng–Refsdal lens of  $M = 0.5$  on top of a gravitational potential with  $\kappa = 0.15$  and  $\gamma = 0.5$ .

We can see the vertical dark stripes at the edges of the map originating from lack of ray hits. Indeed, the shear concentrates the rays in the horizontal direction and expands them in the vertical direction. (This is the case for backward rays. It is the opposite for real photons.) As a result, to fill the full map we should take rays further apart along the horizontal. A factor  $1/(1 - \gamma)$  in the horizontal direction is usually enough to account for this. In this case, the missing rays show up very clearly, but for some lens systems it may not be so obvious. We must therefore always choose a region large enough to prevent flux losses in our calculations. If the shooting region is too large, we are just wasting time by shooting rays that are never going to be collected. In case of doubt, be generous. In general, the lemma to apply is: it is better to waste some rays than to lose flux.

<sup>†</sup> A whole row of pixels composed of  $\sqrt{N_r}$  rows of rays each is also a good choice, although things may get somewhat tricky if  $N_r$  is not a square number.

TABLE 8.4. Code for magmap.py

<pre>import numpy as np import lens as l import matplotlib.pyplot as plt  ny=401 yl=2.  b=np.zeros((ny,ny)) raypix=15.  sqrpix=np.sqrt(raypix) sqrinpix=np.sqrt(1./raypix) ys=2.*yl/(ny-1) xs=ys/sqrpix xl=2.*yl  nx=np.round(2*xl/xs)+1 yr=np.arange(,nx) y, x=np.mgrid[0.0:1.0,0:nx] perc0=5. perc=5. for i in yr:     if ((i*100/nx)&gt;=perc):         perc=perc+perc0         print round(i*100/nx),"% "  x1=-xl+y*xs x2=-xl+x*xs y1,y2=l.TwoPoints(x1,x2,-0.75,0.,0.75,0.,0.5,2.5)  i1=(y1+y1)/ys i2=(72+y1)/ys i1=np.round(i1) i2=np.round(i2) ind=(i1&gt;=0) &amp; (i1&lt;ny) &amp; (i2&gt;=0) &amp; (i2&lt;ny)     i1n=i1[ind]     i2n=i2[ind]     for i in xrange(np.size(i1n)):         b[i2n[i],i1n[i]]+=1         y=y+1.0 b=b/raypix print np.mean(b) plt.imshow(b,vmin=0,vmax=15) plt.show()</pre>	<pre># This is the number of rays per pixel in absence # of lensing. # Rays per pixel square root (rays/pix in one dir)  # Pixel size on source plane # Side of the square area transported back by a ray. # Size of the shooting region at the mage plane # BEWARE. This may need to be larger for # certain lens models!!!! # Number of rays on a column/row at the image plane # This is an array with pixels on y direction # Grid with pixel coordinates for a row at the image # Percentage step for printing progress # Initial value for perc # Loop over rows or rays # Check if we have already completed perc. # Increase perc. # Print progress # Convert pixels to coordinates in the image plane  # Deflect rays. # We can set another lens model here easily. # Convert coords to pixels at the source plane  # Make pixel corrrds integer number  # Indices of rays falling into our source plane # Coordinates of pixels hitting our source plane  # Loop over hits "on target" # Increase magnification at those pixels # Increase the y coordinate of the pixel/rays # Normalize magnification with N_r # Print mean magnification # Show image</pre>
--	--

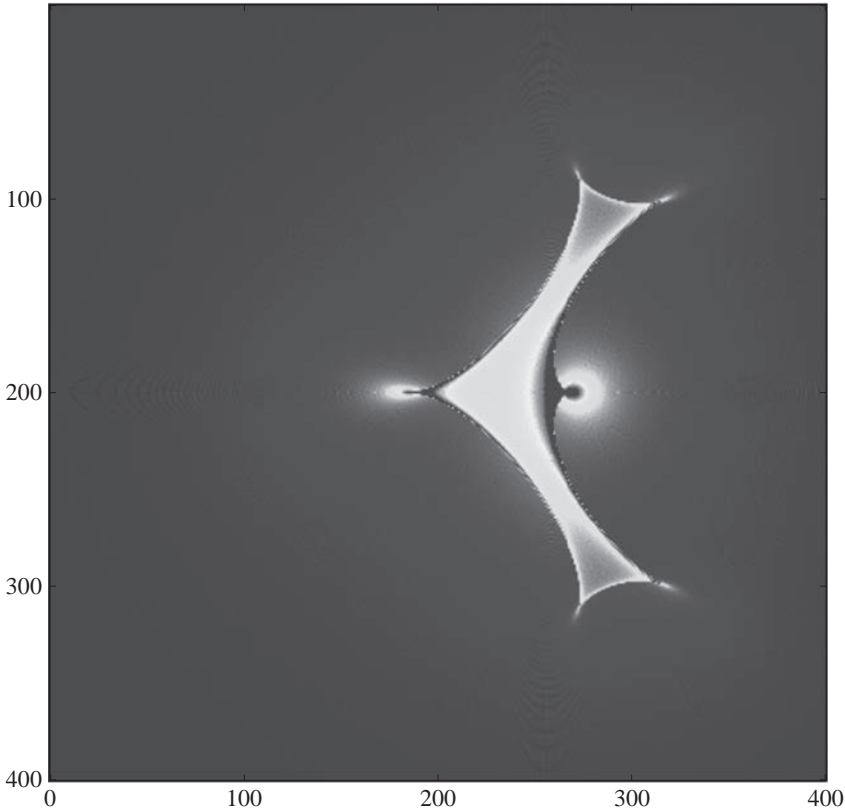


FIGURE 8.5. Magnification map produced with the code **magmap.py**. The lens is a binary with masses 0.5 and 2.5 located at  $(-0.75, 0.0)$  and  $(0.75, 0.0)$ .

#### 8.5.1 Case study: light curves for the binary lens

Magnification maps are, of course, not observable quantities. The source sits at a given location at the source plane and is affected by the magnification at that particular location.<sup>†</sup> But the lens–source system is rarely static, and therefore the magnification with which we see the source changes with time depending on the relative motions of source and lens. The time scale for the change is  $t_E = R_E/v_\perp$  and is called the lensing time or Einstein time (where  $R_E$  is the Einstein radius and  $v_\perp$  is the transverse velocity of the lens with respect to the source–observer line of sight). In galactic microlensing the involved time scales are of the order of tens of days, which is not too long, and it is therefore relatively easy to observe the change in brightness of the source as this relative motion makes the source *travel* through the magnification map. As we have already said, such a brightness vs. time curve is what is called a light curve. A light curve is, therefore, a one-dimensional *cut* through a magnification map. Usually, the cut is performed over a straight line, although in the most general case the path can be curved (e.g. due to parallax).

For quasar microlensing (lensing produced by stars in a galaxy lens that is itself producing multiple macro-images of a background quasar) this time scale is of the order of years. But multiple lens systems have sharp caustics that present large magnification

<sup>†</sup> In fact, if the source has a finite size covering several pixels in our magnification map, it sees an *effective* magnification averaged over its extent.

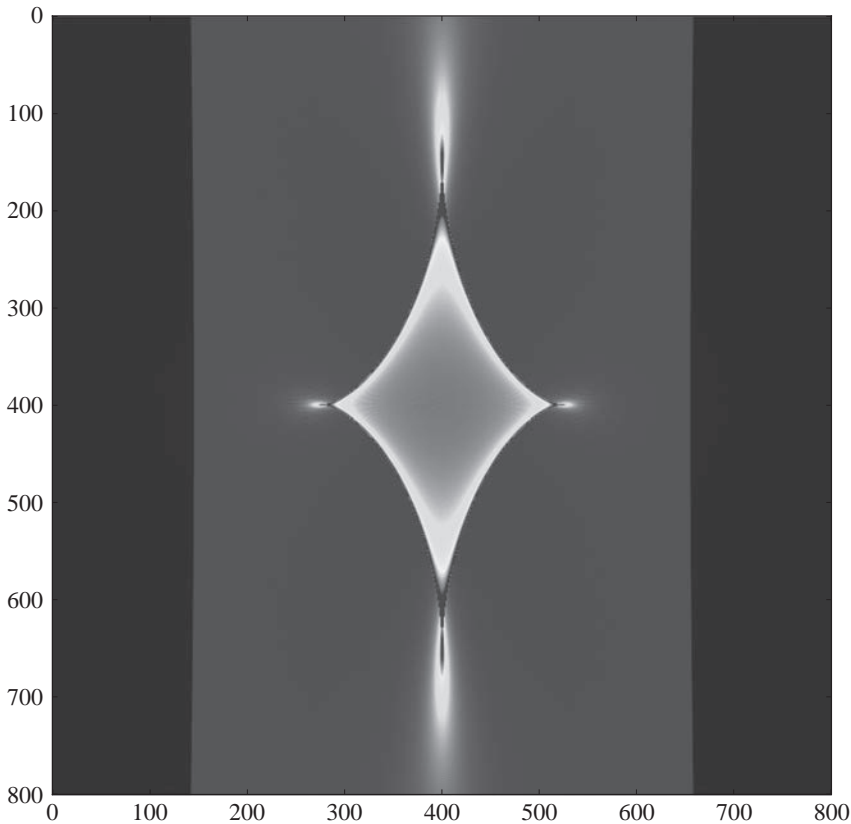


FIGURE 8.6. Magnification map produced with the above-mentioned code. The lens is a Chang–Refsdal lens with mass 0.5 and background potential characterized by  $\kappa = 0.15$  and  $\gamma = 0.5$ .

variations in very narrow regions. Therefore, we can still expect to detect brightness variations on smaller time scales if a source crosses a caustic. To account for this fact, a second time scale called the crossing time is defined as the time it takes a source to travel its own radius (Schneider, Kochanek and Wambsganss 2006). The crossing time is of order of months for typical quasar microlensing systems.

The magnification maps for single axisymmetric lenses are also axisymmetric, and therefore light curves for these systems are symmetric around the time of maximum magnification and, despite providing very valuable information, we could say that they are a bit *boring*. But we also know that a binary lens is enough to spring up a plethora of variety in magnification maps (Schneider and Weiss 1986). Therefore, light curves for binary lenses, being cuts through those maps along a straight line, combine the already rich variety in maps with the variety in paths over the map. The result is an extraordinary richness in light curves, which makes them indeed much more appealing. Light curves for binary lenses are widely used in galactic microlensing for modelling some peculiar events.

As an exercise we will try here to reproduce some of the light curves from binary lenses from a well-known paper by the MACHO Collaboration (Alcock et al. 2000). The paper provides the lens parameters, so we can easily produce the magnification maps. Here, the angular scales will be given in units of the Einstein radius for the total mass of the lens, so we can assume  $M_T = M_1 + M_2 = 1$  (for a different total mass, lengths scale as  $L \propto M_T^{1/2}$ ). The relevant parameters are therefore the mass ratio  $\mu = M_1/M_2$  and lens



separation  $a$ , which we can get from the paper. Thus, two masses  $M_1 = 1/(1 + \mu)$  and  $M_2 = \mu/(1 + \mu)$  located at  $(M_2 a, 0.0)$  and  $(-M_1 a, 0.0)$  are used for the magnification map.

The path along the map is described by the angle with respect to the line between the lenses,  $\theta$ , and the closest approach of the track to the centre of mass of the system,  $u_{\min}$ . This is a straight line  $y = mx + x$  with  $m = \tan \theta$  and  $n = u_{\min}/\sqrt{m^2 + 1}$ . Once the abscissae of the extremes of the track are chosen (for example, as the edges of the magnification map), ordinates can easily be calculated from the straight line equation above. Some other parameters (the time of closest approach to centre of mass  $t_0$ , and a lensing time  $\hat{t}$ ) are needed to put the  $x$ -axis into an absolute time scale. In our case, since we are interested just in the procedure to produce the curves, we will keep it in angular units.

In order to be able to produce the profile, we will use a simple function to produce the cut through magnification map. Let us put it in a module called **aux.py**. The function, which we call **profile**, receives as parameters the magnification map, and the abscissae and ordinates of the initial and end points of the track. The module would look something like this:

```
import numpy as np
from math import sqrt

def profile(c,x0,y0,x1,y1,method='nn'): # Coords are in pixels
    num=int(round(sqrt((x1-x0)**2+(y1-y0)**2))) # Length of track in pixels
    xp, yp = np.linspace(xpp0, xpp1, num), np.linspace(ypp0, ypp1, num)
    # x and y coordinates of track
    zp =c[yp.astype(np.int), xp.astype(np.int)]
    return zp
```

Be aware that **profile** expects coordinates in pixels, not in  $R_E$  units, so we must covert coordinates into pixels first.

With our previous code to produce magnification maps and this function, it is easy to try to reproduce some of the cases in Alcock et al. (2000). Here, we show the result for events 98-SMC-1 in Figure 8.7 and 403-C in Figure 8.8.

The light curves are pretty noisy in the low magnification regions because we have only used 100 rays per unlensed pixel in the calculations.

### 8.5.2 Quasar microlensing magnification maps

Our next step is to address the calculation of magnification maps for the case of quasar microlensing. In quasar microlensing we have a *macro* lens responsible for the multiple images of the quasar. From the image configuration we can calculate a *macro* model, which provides values of the convergence  $\kappa$  and shear  $\gamma$  at the position of each image. But if a fraction  $\alpha$  of the surface mass density originating the multiple images is in the form of compact objects, then microlensing can occur owing to the line of sight to the source passing close to one of these objects. The total convergence  $\kappa$  is therefore split into a part coming from a smooth distribution of matter  $\kappa_s = (1 - \alpha)\kappa$  and a part coming from compact objects (able to produce microlensing)  $\kappa_* = \alpha\kappa$ . This last fraction of matter is assumed here to be in the form of point sources (although other lens models are also possible) and distributed uniformly.

The lens equation for quasar microlensing is then:

$$\mathbf{y} = \begin{pmatrix} 1 - \kappa_s - \gamma & 0 \\ 0 & 1 - \kappa_s + \gamma \end{pmatrix} \mathbf{x} + \sum_{i=1}^{N_*} m_i \frac{(\mathbf{x} - \mathbf{x}_i)}{|\mathbf{x} - \mathbf{x}_i|^2}. \quad (8.9)$$

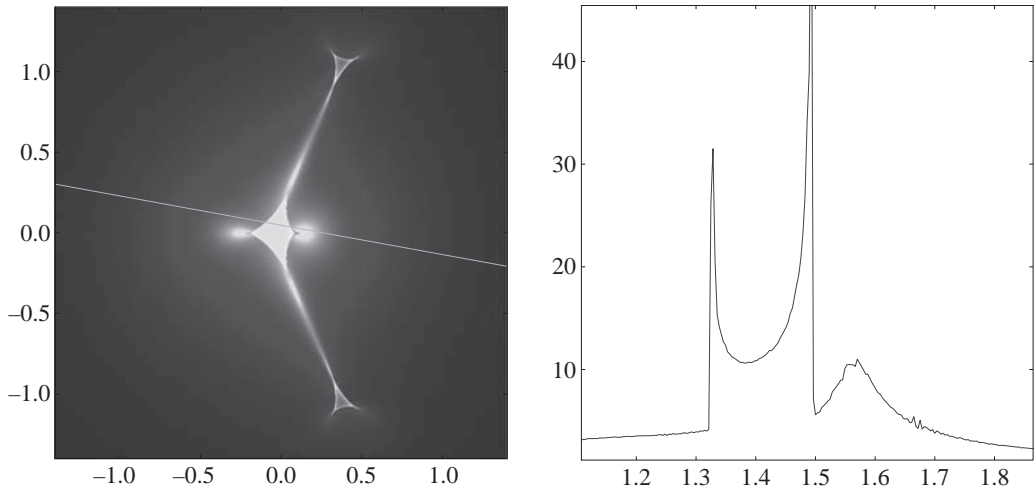


FIGURE 8.7. Magnification map, source track and light curve for microlensing event MACHO 98-SMC-1.

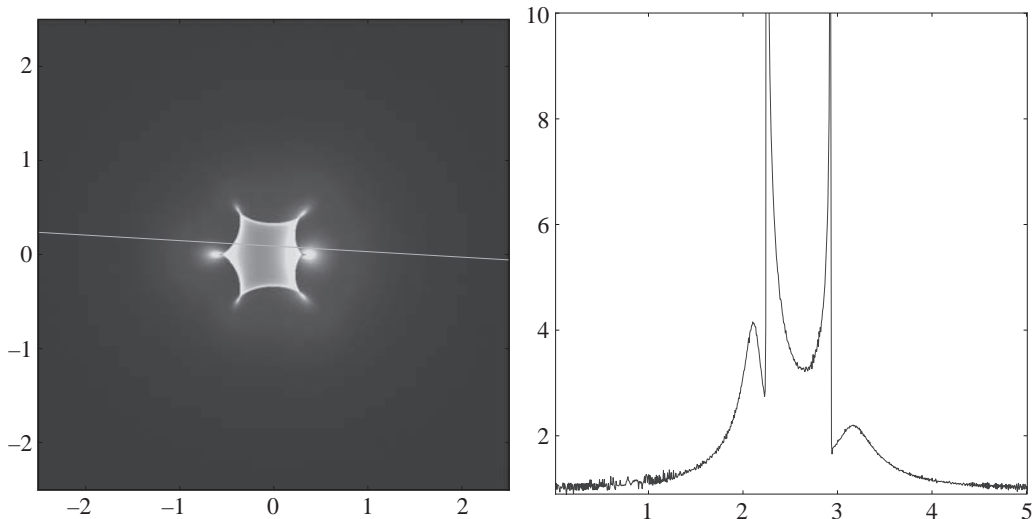


FIGURE 8.8. Magnification map, source track and light curve for microlensing event MACHO 403-C.

For simplicity we assume here that all masses are equal to one solar mass  $m_i = m = 1$ , although different mass distributions can be considered as well. As we are making our calculations in units of the Einstein radius for a unitary mass, the effect of having a different average mass is just to rescale all length scales by  $\sqrt{\langle m \rangle}$ .

We choose a rectangular shooting region at the lens plane with axis ratio  $q = (1 - \kappa + \gamma)/(1 - \kappa - \gamma)$ . The reason for this is that this is the axis ratio of the region into which a square at the source plane would be mapped by the lens if all the convergence were to be in smooth matter. The size of the shooting region is made proportional to the size of the magnification map, taking into account the stretching mentioned above. An extra factor  $3/2$  is included to account for the fact that the graininess of the deflector may deflect rays

further out into our magnification map. Therefore, if we want a magnification map with half-size  $y_l$ , the shooting region will have a size  $1.5y_l/|(1 - \kappa - \gamma)|$  and  $1.5y_l/|(1 - \kappa + \gamma)|$  along the horizontal and vertical axes respectively.

Another important detail to be considered is the size of the region where we are going to distribute the deflectors. The issue here is the *diffuse flux* that may end up hitting our map at the source plane due to deflectors outside the region where we are putting our stars. Katz, Balbus and Paczynski (1986) and Schneider and Weiss (1987) independently estimated that the minimum number of stars that need to be included in order to lose less than a fraction  $\epsilon$  of the flux is given by

$$N_* = \frac{3\kappa_*^2\epsilon^{-1}}{(1 - \kappa)^2 - \gamma^2}. \quad (8.10)$$

Taking into account that for a square region of side  $L_x$  we have

$$N_* = \kappa_* L_x^2 / \pi,^\dagger \quad (8.11)$$

this gives a minimum size for the region into which the deflectors are to be distributed of

$$L_x^{\min} = \left( \frac{\pi 3\kappa_* \epsilon^{-1}}{(1 - \kappa)^2 - \gamma^2} \right)^{1/2}. \quad (8.12)$$

This size is valid to ensure that there is no more than  $\epsilon$  flux lost at the origin. But the shooting region may be quite large for large magnification maps (and/or large values of the mean magnification). This size then has to be increased accordingly (and also, of course, the number of stars in order to fulfil equation 8.11) by the size of the shooting region ( $2x_l$ ):

$$L_x = L_x^{\min} + \frac{3y_l}{\min(|1 \pm \kappa - \gamma|)}, \quad (8.13)$$

where ‘min’ in the denominator of the second term chooses the sign in  $\pm$  that makes the second term larger.

For most practical applications  $\kappa_*$  is a very small number and (except, for example, for very small sizes of the magnification maps) the second member of equation (8.13) usually dominates, and it is usually safe to neglect the first term. In any case, it is necessary to compare both terms before neglecting either of them. In our implementation, once this size is determined, the number of stars is calculated using equation (8.11).

We are now ready to write a piece of **Python** code that calculates magnification maps for quasar microlensing. The algorithm will be structured the following way:

- (i) Set up the input parameters of the problem: total convergence ( $\kappa$ ) and shear ( $\gamma$ ), the fraction of convergence in the form of stars  $\alpha$ , rays per unlensed pixel ( $N_r$ ), the size of the magnification map in Einstein radii ( $2y_l$ ) and the resolution as the number of pixels of a side of the map ( $N_y$ ). We may also set the value for the fraction of flux  $\epsilon$  that we may leave out, and that will set the minimum number of stars and the size of the region populated with deflectors. We then make some calculations to set up the size of the shooting region ( $2x_{l1}, 2x_{l2}$ ) and of the region populated with the stars ( $2x_{ls}$ ) according to the principles described above. Finally, in this preparatory section, we create the array that is to contain our final map.

<sup>†</sup> It is sometimes common to distribute the stars in a circle. In that case, if  $L_x$  is the radius of the circle, the  $\pi$  factor in the denominator disappears.

- (ii) Next, we randomly distribute  $N_*$  stars within the appropriate region.
- (iii) Loop over the grid of rays in a similar way to what we did in Section 8.5. Again, we benefit from **numpy** arrays and throw one row of rays at a time in order to speed up the process.
- (iv) Deflect the row of rays according to the lens equation. This deflection contains an inner loop over all the deflectors which, for many interesting cases, is quite time consuming.
- (v) With the coordinates of the deflected rays at the source plane, calculate the coordinates of the pixel on which the rays have hit and, if it is within our region of interest, add 1 to the value of that pixel.
- (vi) When the loop over rows of rays is over, normalize the magnification map by dividing the array by the number of rays per unlensed pixel.
- (vii) Display and/or save the results.

A **Python** code that does just that is shown in Table 8.5. This program is somewhat longer than our previous examples, but still it is only 131 lines of code, of which nearly 20% are mostly *ornamental* (i.e. printing values and comments). It is pretty simple yet it is a fully stand-alone piece of software (assuming the needed libraries are installed, of course) capable of producing nice magnification maps for quasar microlensing. An example of a magnification map produced with this code can be found in Figure 8.9. On top, the lens plane is also shown with the position of the stars and a box marking the shooting region.

Figure 8.10 shows the magnification map for a real lensed quasar. It is a magnification map corresponding to the image D of the famous Einstein Cross lens system QSO 2237+0305. This lens system is peculiar in many ways, but particularly in that the lines of sight through the lens that produce the four images of the quasar cross the very innermost part of the lens galaxy; therefore, we can assume that nearly 100% of the surface mass density is in the form of stars. But for most multiply imaged quasars, this is usually not the case. Mediavilla et al. (2009), by studying a sample of 29 image pairs through 20 lens systems, showed that, on average, around 5% of the mass is in the form of stars while the rest is smooth (dark) matter.

## 8.6. Source size effects

Up to now, we have implicitly assumed that the source has the size of one pixel at the source plane. We know nothing about the magnification structure at subpixel scales. The only solution is to increase the resolution of the magnification map. On the other hand, if the source is larger than a pixel, the source is affected by an *effective* magnification, which is the convolution of the source's brightness distribution with the magnification map at that location. If the source has a brightness distribution described by  $S(\mathbf{y})$ , then the magnification of a finite source  $M^s$  at location  $\mathbf{y}$  at the source plane would be given by

$$M^s(\mathbf{y}) = \sum_{\mathbf{y}_0} S(\mathbf{y} - \mathbf{y}_0) M(\mathbf{y}_0), \quad (8.14)$$

where  $M(\mathbf{y}_0)$  is the magnification at pixel  $\mathbf{y}_0$ , and the summation is carried over the whole magnification map (although only pixels within the extent of the source do actually contribute). Therefore, the effect of a finite source size (larger than the pixel size) is to smear the magnification map structures. The larger the size of the source, the larger the smearing effect. The effect of different source sizes is illustrated in Figure 8.11.

TABLE 8.5. Code for qmic.py

```

import numpy as np
from math import pi
import matplotlib.pyplot as plt
from random import seed, uniform
from time import time, clock, sleep
from pyfits import writeto
startt=time()
# ***** Model Parameters *****
kappa=0.59
gamma=0.61
alpha=0.999
raypix=15.0
ny=1000
yl=10
eps=0.02
# **** Make some preliminary calculations*****
ks=kappa*alpha
kc=kappa*(1.-alpha)
ys=2.*yl/(ny-1)
ooy=1./ys
sqrpx=np.sqrt(raypix)
f1=1./abs(1.-kappa-gamma)
f2=1./abs(1.-kappa+gamma)
fmax=max(f1,f2)
xl1,xl2=1.5*yl*f1, 1.5*yl*f2
xl=1.5*yl*fmax
nsmin=3*ks**2/eps/abs((1.-kappa)**-gamma**2)
xmin=np.sqrt(pi*nsmin/ks)/2
xls=xl+xmin
nx1=np.int32(np.round(1.5*ny*f1*sqrpx))
nx2=np.int32(np.round(1.5*ny*f2*sqrpx))
nx=max(nx1,nx2)
xs=2.*xl1/(nx1-1)
xnl=abs(ks*(2*xls)*(2*xls)/pi)
nl=int(xnl)
thmag=1./(1-kappa-gamma)/(1-kappa+gamma)
print "*****"
print "Half size of map in Einstein radii =", yl
print "Number of pixels of magnification map =", ny
print "Half size of shooting region =", xl
print "Number of rays along the longest axis =", nx
print "Half size of region with microlenses =", xls
print "Total convergence k =", kappa
print "Shear gamma =", gamma
print "Fraction of mass in microlenses, alpha =", alpha
print "Convergence in form of microlenses, ks =", ks
print "Number of microlenses =", nl
print "Rays per unlensed pixel, r aypix =", raypix
print "Theoretical mean magnification mu =", thmag
print "*****"
b=np.zeros((ny,ny))
# ***** Randomly distribute stars in region *****
x1l=np.zeros(nl)
x2l=np.zeros(nl)
seed(1.0)
for i in range(nl):
    x1l[i]=uniform(-xls,xls)
    x2l[i]=uniform(-xls,xls)
print "*****"
perc0=5.
perc=5.
yr=np.arange(0,nx2)
y,x=np.mgrid[0.0:1.0,0:n1]

nlrange=np.arange(nl)
# ***** MAIN LOOP *****
for i in yr:
    if ((i*100/nx2)>=perc):
        perc=perc+perc0
        print round(i*100/nx2),"% ", round(time()-startt,3), " secs"
        # print completed fraction and elapsed execution time
        x2=-x2l+y*xs
        x1=-x1l+x*xs

```

```

# Import needed modules

# Random number stuff
# Timing stuff
# To be able to save output as fits
# Time at start of execution

# Total Convergence
# Shear
# Fraction of mass in form of microlenses
# Rays per pixel in absence of lensing
# Pixels in the magnification map
# Half size of magnification map in Einstein radii
# Maximum fraction of flux lost

# Convergence in microlenses
# Convergence in smooth matter
# Pix size in the image plane
# Inverse of pixel size on image plane
# Rays per pixel in one dimension
# Exp. factor on horizontal axis
# Exp. factor on vertical axis
# Max Exp. factor
# Half size of shooting region in x and y
# Longest half side of shooting region
# Min number of stars
# Min half side of star region
# Expand to account for shooting region
# Rays in shoot. reg. along x axis
# Rays in shoot. reg. along y axis
# Number of rays along longest side
# Pixel size on image plane
# Number of microlenses
# Number of microlenses (int)
# Theoretical value of magnification

# Print some parameters

# Initialize magnification map

# Initialize microlens positions to zero

# Initialize random number generator
# Generate positions of microlenses

# Percentage step to show progress
# Initial percentage
# Array for looping over rows of rays
# These are arrays with x and y coords
# of one row of rays in image plane
# Array for looping over lenses

# Main loop (over rows of rays)
# If perc is completed, then show progress

# Convert pixels to coordinates in the image plane

```

(cont.)

TABLE 8.5. (*cont.*)

y2=x*0.0	# Initialize variables
y1=x*0.0	
for ii in nrange:	# Loop over microlenses
x1ml=x1-x1l[ii]	
x2ml=x2-x2l[ii]	
d=x1ml**2+x2ml**2	# Distance to lens ii squared
y1=y1+x1ml/d	# Deflect x coordinate due to lens ii
y2=y2+x2ml/d	# Deflect y coordinate due to lens ii
del x1ml,x2ml,d	
y2=x2-y2-(kc-gamma)*x2	# Calculate total y deflection
y1=x1-y1-(kc+gamma)*x1	# Calculate total x deflection
i1=(y1+y1)*oofs	# Convert coordinates to pixels on source plane
i2=(y2+y1)*oofs	
i1=np.round(i1)	# Make indices integer
i2=np.round(i2)	
ind=(i1>=0) & (i1<ny) & (i2>=0) & (i2<ny)	# Select indices of rays
	# falling onto our source plane
i1n=i1[ind]	# Array of x coordinates of rays within map
i2n=i2[ind]	# Array of y coordinates of rays within map
for ii in xrange(np.size(i1n)):	# Loop over rays hitting the source plane
b[i2n[ii],i1n[ii]]+=1	# Increase map in one unit if ray hit
y=y+1.0	# Move on to nex row of rays
print "*****"	
b=b/raypix	# Normalize by rays per unlensed pixel
print "*****"	
print "Measured mean magnification =",np.mean(b)	
print "Theoretical magnification is =",thmag	
print "*****"	
if (thmag<0):	# Vertical or horizontal flip in some cases
if (gamma<0):	
b=np.flipud(b)	
else:	
b=np.fliplr(b)	
# ***** Display results *****	
ax=plt.subplot(121)	# Left plot
plt.plot(x1l,x2l,'+')	# Plot positions of stars
rayboxx=[-x1l,-x1l,x1l,x1l,-x1l]	
rayboxy=[-x2l,x2l,x2l,-x2l,-x2l]	
plt.plot(rayboxx,rayboxy)	# Show shooting region
mapboxx=np.array([-y1,-y1,y1,y1,-y1])	
mapboxy=np.array([-y1,y1,y1,-y1,-y1])	
plt.plot(mapboxx*f1,mapboxy*f2,'r')	# Show region mapped onto map
plt.xlim(-1.1*xls,1.1*xls)	
plt.ylim(-1.1*xls,1.1*xls)	
ax.set_aspect('equal')	# Keep aspect ratio
plt.subplot(122)	# Right plot
imshow=plt.imshow(b,origin='lower')	# Display magnification map
print "*****"	
print "Exec. time = ",round(time()-startt,3), ' seconds'	# Print execution time
plt.show()	
# ***** Save result as fits file? *****	
save=''	
while (save not in ['y','n']):	# Wait for input unless it is 'y' or 'n'
save=raw_input("Save file (y/n)? ")	
if (save == 'y'):	# If 'y'
filename=''	
filename=raw_input("Filename = ")+'fits'	
writeto(filename,b)	# Write fits file

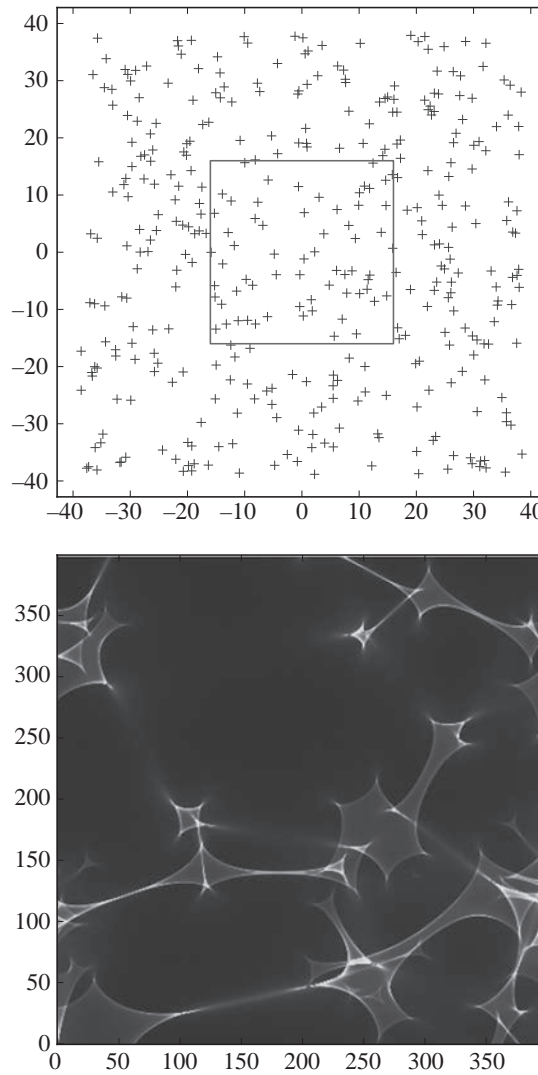


FIGURE 8.9. Magnification map produced with the **qmic-pv.py** code for  $\kappa = 0.7$  and  $\gamma = 0.0$ . 28.6% of the lens mass is in the form of stars, resulting in  $\kappa_* = 0.2$ . The size of the map is  $6.4R_E$  and it has  $400 \times 400$  pixels. Top: the lens plane with the positions of the 385 stars. The ray-shooting region is marked by the square.  $\epsilon$  is set to 0.01.

We may wonder to what extent the effect of the size depends on the shape of the luminosity profile of the source. Mortonson, Schechter and Wambsganss (2005) have shown that for circular sources the radial profile of the source is quite irrelevant for statistical purposes. The only important parameter is the *half-light radius*  $r_{1/2}$  that contains half of the source luminosity. Once we know that the particular shape of the brightness profile of the source is not that important, we are free to choose the handiest one. It is quite common to use a circular Gaussian of sigma  $r_s$  to model the source. The half-light radius for such a source is  $r_{1/2} = 1.18r_s$ . Other profiles have also been used in the literature (e.g. uniform discs, etc.). Gaussians are very convenient for modelling the source, as convolution with a Gaussian is a pretty common task, and it is very easy to find the



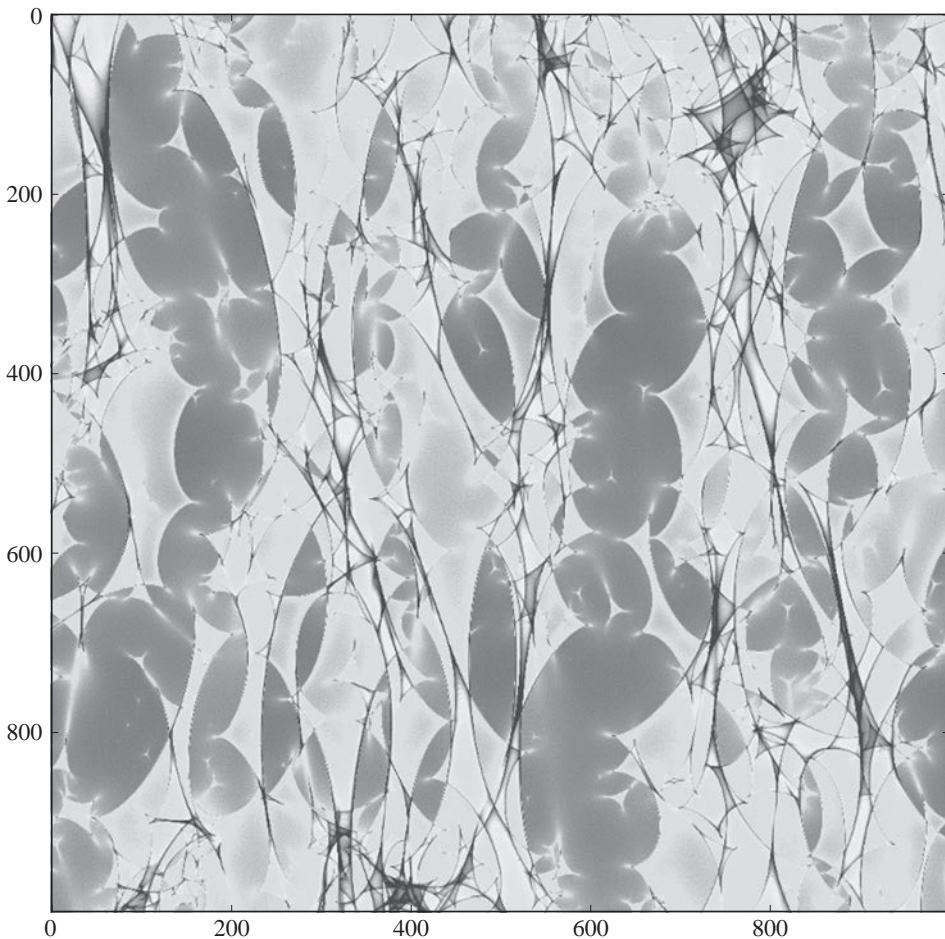


FIGURE 8.10. Magnification map for image D of QSO 2237+0305. The map is  $20R_E$  on a side. The parameters for the calculations are  $\kappa = 0.59$ ,  $\gamma = 0.61$  and 100% of the mass in the form of stars. The calculation was done with 6553 stars.

convolution already implemented in most programming languages.<sup>†</sup> **Python** is no exception, and we can find such a function in the **ndimage** module of **scipy**. The function is called **gaussian\_filter**. We can therefore add a function for the Gaussian convolution function **gaussconv** into **aux.py** to perform this task:

```
import scipy.ndimage # To go at the top of aux.py with the other imports
def gausscon(a,size): # a is the input map and size the sigma of the gaussian
    b=a*0.0          # Create the output array. Same size as input array
    scipy.ndimage.gaussian_filter(a, [size, size],output=b)
    return b
```

The effect of this blurring can be seen more clearly on the light curves. A cut through those maps can be seen in Figure 8.12. We clearly see that sharp changes in magnification

<sup>†</sup> 2D convolution with a Gaussian kernel is a separable problem that can be split into two 1D convolutions. Implementations that benefit from this are much faster, particularly for large kernels.

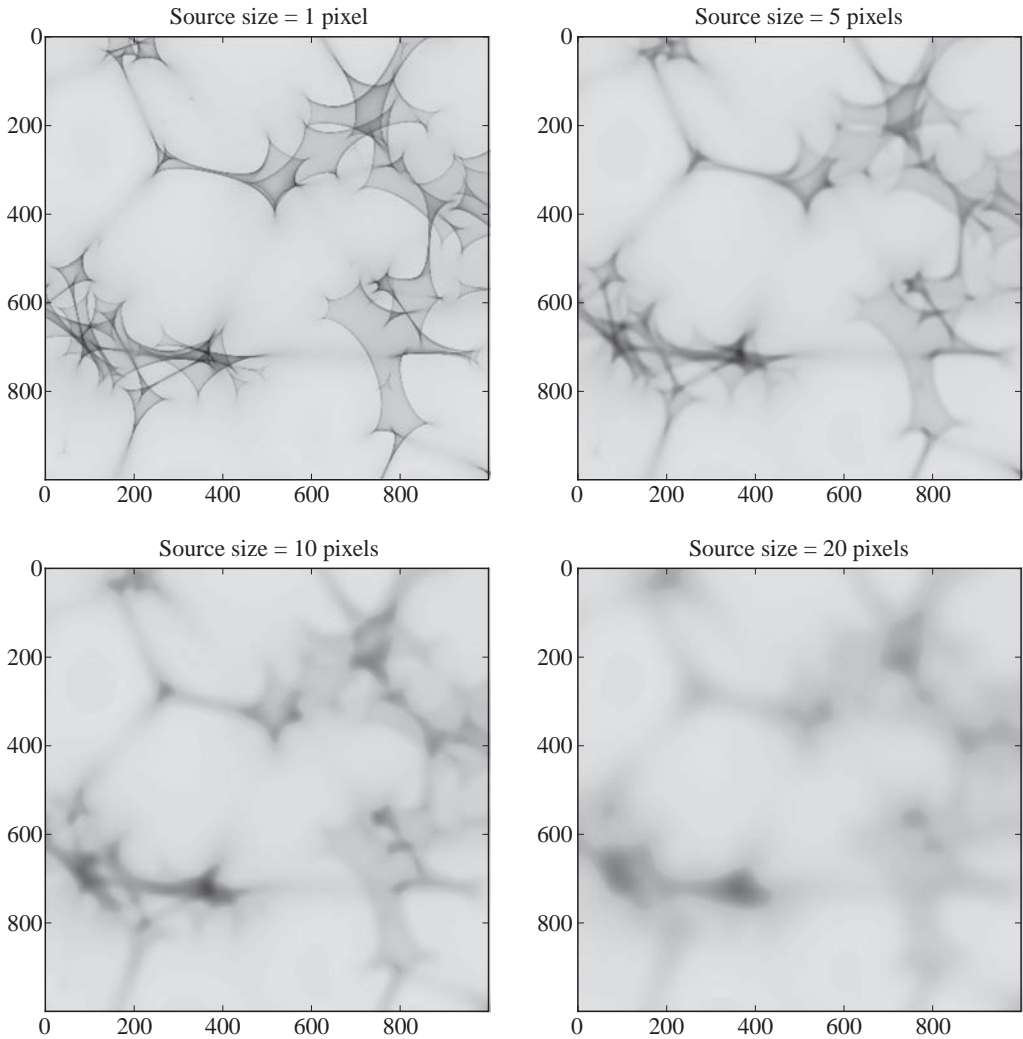


FIGURE 8.11. The effect of source size on effective magnification. The original map (upper left) corresponds to  $\kappa = 0.45$  and  $\gamma = 0$  with 90% of the matter in the form of stars of mass  $M = 1 M_{\odot}$ . The side of the map is  $10 R_E$ . For the other maps, the source has a Gaussian profile with a  $\sigma$  of 5 (upper right), 10 (lower left) and 20 (lower right) pixels (corresponding to 0.005, 0.01 and 0.02  $R_E$  respectively).

due to caustic crossings disappear for large sources, and that light curves look pretty smooth for the larger sizes. These different light curves for the same lens system but for different sizes are indeed to be expected. The standard model for accretion discs around black holes is that of Shakura and Sunyaev (1973), who predicted that the temperature of the accretion disc is much higher in the inner parts than further out. Thus, it also predicts that the *effective* radius of the disc is wavelength dependent (with a power law  $r \propto \lambda^{4/3}$ ) and we can expect to be able to observe something like the light curves of Figure 8.12 for different wavelengths. For example, the blue curve could correspond to X-ray observations while the cyan one could correspond to the optical disc. This effect is known as microlensing *chromaticity*. A recent study of this effect using these techniques on *HST* data has been done by Muñoz et al. (2011).

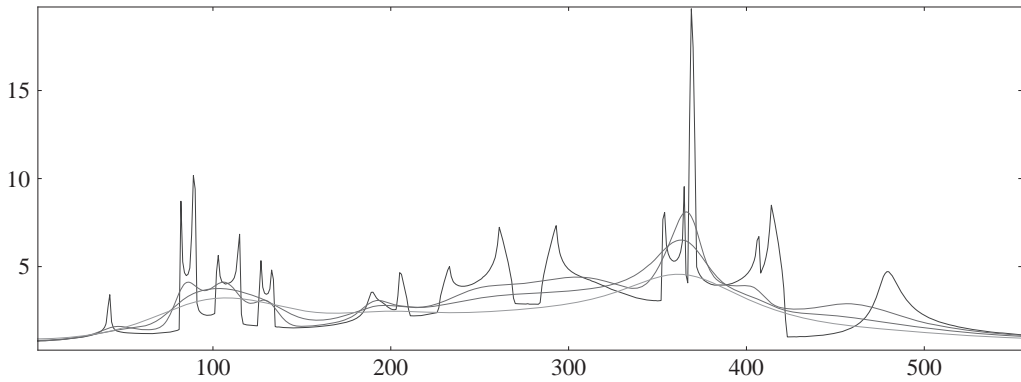


FIGURE 8.12. Light curve for different sizes of the source. Blue, green, red and cyan correspond to Gaussian sources with a sigma of 1, 5, 10 and 20 pixels respectively shown in Figure 8.11.

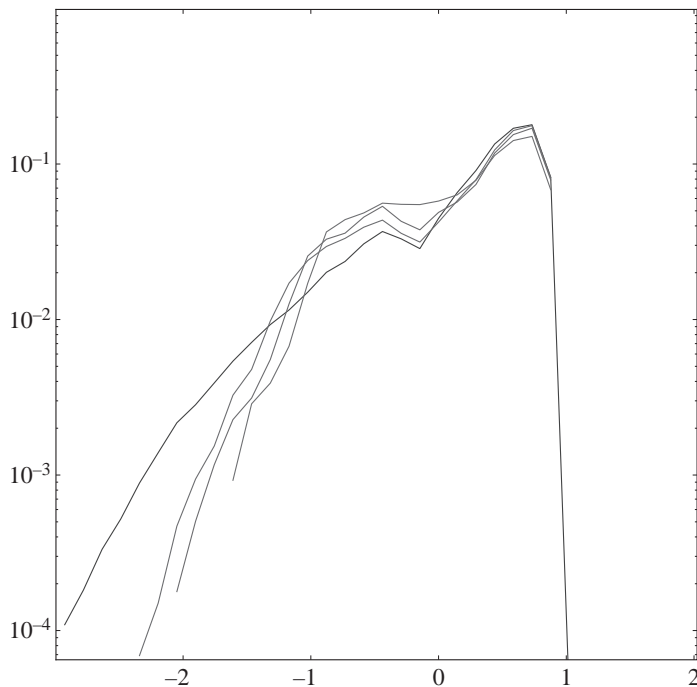


FIGURE 8.13. Magnification histograms for different sizes of the source. Blue, red, green and magenta correspond to Gaussian sources with a sigma of 1, 5, 10 and 20 pixels respectively shown in Figure 8.11.

To see what the effect of this blurring is on the statistics of the magnification maps, we may plot histograms of the magnification in the different maps shown in Figure 8.11. As astronomers are used to expressing brightness in magnitudes, it is customary also to represent magnification histograms in magnitudes around the mean magnification of the map. The result is shown in Figure 8.13.

What we see is that the larger the size, the less probable it is to find large magnifications. The rule of thumb is therefore: *small sources suffer more microlensing magnification than large sources*. Of course, the latter is to be understood in a statistical sense. But beware: Even things with very low probability may eventually happen! So do not

take the above as an infallible rule. In the particular case shown in Figure 8.13, this effect is more pronounced for negative magnifications. So, it is nearly 8 times more probable to find a magnification of  $-2$  mag for a source of 1 pixel (blue curve) than it is to find it for sources of 10 pixels (green curve). To the contrary, lower magnifications below  $-1$  are more probable for larger sizes.

This effect therefore allows us to perform statistical studies based on the observed microlensing magnifications and to use this effect to put constraints on the size of the source. This has been done for example by Jiménez-Vicente et al. (2012) to estimate the size of the accretion disc for a sample of 19 lensed quasars.

## 8.7. Beyond simple inverse ray shooting: treecodes and IPM

Inverse ray shooting is a pretty simple algorithm with a very easy implementation, as we have just seen. Nevertheless, for many practical applications it requires very long execution times. Quite often we need to address problems in which very compact and very large sources are involved at the same time. This means that we need to calculate magnification maps for very large regions at the source plane but, at the same time, with a very high resolution. The reason why these calculations take so long is easily found by looking at equation (8.8). As an example, let us work out how many calculations are needed to produce a map at moderately high resolution (i.e.  $N_x = 2500$ ) and for a large region at the source plane (which means we will need to include many stars at the lens plane in our calculations). Let us take  $N_* = 5 \times 10^5$ . To have good accuracy in the map we need to shoot a few hundred rays per unlensed pixel. Let us take  $N_r = 500$ . If every deflection needs of order 10 operations, we have, for the total number of operations to deflect all the rays:

$$N_{\text{op}} = 10 \times N_x^2 \times N_r \times N_* \approx 1.5 \times 10^{16} \text{ flop} \approx 15 \text{ Pflop}. \quad (8.15)$$

This is very heavy indeed. A modern fast workstation is able to sustain around 10 Gflop per second. Such a computer would need nearly three weeks to finish this tremendous task! Even a much more modest calculation with only  $N_x = 1000$  pixels and  $N_* = 5 \times 10^4$  stars would need nearly 7 hours!

Of course, we can always use faster hardware. But not everyone may have access to it. Moreover, using it may sometimes be rather complicated. Nowadays, the use of GPUs for scientific computing has become increasingly popular. They provide very fast yet reasonably cheap hardware that is perfectly suited to the calculation of microlensing magnification maps. This is because GPUs benefit from having many available cores when the problem is highly parallelizable. And it is difficult to find a problem more parallelizable than this one! Indeed, every ray deflection is independent of any other, and even the deflection produced by a lens/cell is independent of the others. This makes the load balance very easy, and this problem very well suited for GPUs. A performance comparison of an implementation of a plain IRS algorithm into a GPU with a parallel implementation of the hierarchical treecode into a cluster of CPUs and a standard implementation of the hierarchical treecode into a single CPU can be found in Bate et al. (2010).

How can we help these calculations go faster? A careful look into the previous equation reveals that there are only two places where there is some hope, if any, of making improvements. We cannot change the resolution of the map as this is set by the problem. We may try to optimize the number of operations needed to perform a single-ray deflection. But, to be honest, even if we do the finest of jobs, there is not much gain. But what about the other two factors? The reader may (rightly) say that the number

of stars is also set by the problem, and that a high number of rays per unlensed pixel is needed in order to produce accurate maps. We may fine tune the latter number a bit, but, again, to no avail. But some clever people have seen somewhat deeper than that, and realized that we can indeed save quite a large amount of time by reducing those two factors. The reduction of the number of lenses included in the calculations is the basis of treecodes. The reduction of the number of rays per unlensed pixel needed to produce a highly accurate map is the path taken by the developers of the inverse polygon mapping (IPM) algorithm.

### 8.7.1 Treecodes

The key idea behind treecodes is to group particles together and treat their combined gravitational effect as if it came from a single entity. This is possible because the effect of particles far away is weak (it diminishes as  $1/r$ ) so that many particles can be grouped together and their combined effect taken into account as a single *pseudoparticle*. This way, we do not need to loop over all the lenses for every single ray.

The effect of many of those far away lenses varies quite smoothly with position, and we can then group together a bunch of lenses that are far away from a given point but close to each other and treat them as a single pseudoparticle whose properties very much reproduce the gravitational effect of the whole set of individual lenses. This way, for a given location at the lens plane, we only really need to treat individually the closest lenses. For the rest, they can be grouped together into pseudoparticles. This way, the deflection angle is split into two parts:

$$\alpha = \sum_{i=1}^{N_*} \alpha_i = \sum_{i=1}^{N_*^{\text{near}}} \alpha_i + \sum_{i=1}^{N_*^{\text{pseudo}}} \alpha'_i = \alpha^{\text{near}} + \alpha^{\text{pseudo}}, \quad (8.16)$$

where  $N_*^{\text{near}}$  is the number of nearby lenses that are to be included individually, and  $N_*^{\text{pseudo}}$  is the number of *pseudolenses* (of which each accounts for the contribution of a given number of individual lenses). This way, the factor  $N_*$  in our equation is greatly reduced as, now, the effective number of stars is  $N_* = N_*^{\text{near}} + N_*^{\text{pseudo}}$ , which may be of the order of a few hundred. No wonder there is such a big gain in execution speed by using this approach.

This idea of grouping together faraway lenses to speed up computation of the magnification maps was already used in the seminal work by Schneider and Weiss (1987). The procedure was refined and improved by introducing a *hierarchical* treecode by Wambsganss (1990, 1999). Hierarchical treecodes were introduced into  $n$ -body problems in the early 1980s by several authors, the most popular and widely used being the algorithm introduced by Barnes and Hut (1986).

The idea of a hierarchical treecode is that of recursively dividing the space into smaller cells at different levels. The division ends when every cell has either one or zero particles. It is common that cells at a given level are half the (linear) size of the cells in the immediately inferior level. In 3D, if we start with a cubic volume (the *root* cell) and divide it into eight cells (branches) of half the (linear) size, and we repeat the procedure recursively, this division of space is called an *octree*. In 2D, the equivalent division of space has four subcells per parent cell and the tree is called a *quadtrees*. For each cell, we calculate the properties of a pseudoparticle with the total mass of the particles within the cell and located at the centre of mass. We can even calculate higher multipolar terms to account for finer details of the gravitational effect of the pseudoparticle. The idea behind hierarchical treecodes is to use larger cells for faraway regions, smaller cells for

nearer regions, and individual particles for the nearest region. In principle, we would like to use cells as large as possible without introducing large errors. How do we decide if a cell has the right size to stop going *down* the tree?<sup>†</sup> In the Barnes–Hut algorithm this is decided via the *cell-opening* parameter  $\delta$ . This criterion establishes that a cell of size  $s$  is already small enough to be included in the calculation if it is at distance  $d$  and fulfils that  $s/d < \delta$ . Values of  $\delta \lesssim 1$  usually give a good approximation for the force calculations.

This way, if we are calculating the forces between  $N$  particles, the algorithm reduces a problem from  $O(N^2)$  to  $O(N \log N)$  computations. This is indeed a huge improvement when  $N$  is a large number! In our case, we want to calculate the deflection of  $N_r \times N_x^2$  rays by  $N_*$  stars, so we convert a problem of  $O(N_r \times N_x^2 \times N_*)$  into a problem of  $O((N_r \times N_x^2) \log N_*)$ . Again, a very substantial gain.

A thorough description of the algorithm is completely beyond the scope of this tutorial, and the reader can find all the details in the original publications by Wambsganss (1990, 1999). Nevertheless, I briefly describe here the most important details:

- The lens plane is recursively divided into a quadtree.
- For each cell, not only the centre of mass, but multipoles up to order 6 are calculated.
- An *opening* parameter  $\delta$  with values between 0.4 and 0.9 is used to decide the size of the cells to include in the calculations.
- Further speed-up is achieved by taking into account that the particle/cell structure to include for nearby rays does not change much. Therefore, there is no need to recalculate that structure for every ray, as it can be re-used for a certain amount of them.
- Still further speed-up is achieved by noticing that the deflection angle of faraway cells changes little and very smoothly with position, so it can be calculated at a few positions and interpolated in between.
- The deflection is calculated by directly adding the effect of nearby lenses and including only large cells/pseudoparticles at longer distances.

This way, for problems with millions of particles, speeds increased by up to a factor of a thousand can be achieved.

Using a treecode is not the only way of separating long and short range effects in this context, although it is probably the most widely used. Nevertheless, alternative methods have also been used (e.g. Kochanek 2004 used Fourier transform methods to deal with this issue).

### 8.7.2 Inverse polygon mapping

Further improvement in the performance of the inverse ray-shooting algorithm was introduced by the inverse polygon mapping algorithm (Mediavilla et al. 2006, 2011).

The basic idea behind this algorithm is noticing that what we actually do to calculate the magnification by inverse ray shooting is to transport test areas *backwards* at the lens plane and compare them (by collecting them) with reference areas (pixels) at the source plane. In doing this, we are naively assigning the whole area transported by a ray to the pixel at the source plane that is hit. But this is extremely inefficient, as it ignores a large fraction of information on the lens mapping, and we can actually do much better than that. How? By really transporting a test area backwards and by apportioning it among the pixels that are enclosed by it (if there are more than one). In fact, except for the critical curves, the lens mapping is a diffeomorphism (meaning that the function and its inverse are differentiable), so we can expect it to be topologically *well behaved*. This way, we can use the *backwards* image of a cell in the source plane to calculate the

<sup>†</sup> We should actually say *up* the tree, as we are going from root to leaves. But it is customary to represent the tree upside-down so that going deeper into the tree is going *down*.



local properties of the mapping and, with them, we can make a good estimate of the magnification. What makes this method so good is that the places where we need many rays to have a good estimate of the magnification are those with low magnifications. But those regions are away from the caustics and therefore the backwards mapping of test areas works best even for very large cells! And what about cells containing a critical curve? In principle, these have to be identified and treated separately. Although we will soon see that this is indeed very easily done.

Again, as with the hierarchical treecode, describing all the details of the algorithm here is not possible, and the reader is referred the original papers (Mediavilla et al. 2006, 2011), but we will sketch the main points here:

- (i) Tessellate the lens plane. While the regular grid is a natural choice for the IRS algorithm, any tessellation is possible. Square cells (maybe with an extra ray at the cell centre) are commonly used.
- (ii) Transform cell vertices with the lens equation.
- (iii) Check for out of linearity (i.e. identify critical cells).
- (iv) Reprocess critical cells. This can be done by subdividing them into smaller (hopefully non-critical) cells, applying IRS to these cells, etc.
- (v) Apportion the transformed area among the corresponding pixels at the source plane.

As stated in Mediavilla et al. (2006), the IRS would correspond to steps (i) and (ii) and would be the zeroth-order approximation. IPM without a linearity check would take steps (i), (ii) and (v) and would be first order. IPM with a linearity check would use the five steps and would correspond to a second-order approximation. The good news is that first-order IPM, without any linearity control (i.e. ignoring critical cells and treating them as if they were not critical) works so well that it is quite satisfactory. The gain with IPM comes from the fact that the cells that we are transforming back can be as large as the pixels we used with IRS at the image plane or even larger. That is equivalent to saying that we are using  $N_r \approx 1$  or even smaller! In fact, values of  $N_r$  as small as 0.12 (i.e. transported cells are three pixels wide!) are able to produce magnification maps even more accurate than standard IRS, with 350 rays per pixel, in a tiny fraction of the computational time. Indeed, moving from  $N_r = 500$  to  $N_r = 0.2$  in equation (8.8) makes a huge difference. With this procedure, if done carefully, the extra time to set up the tiling, etc., is really minimal and accounts for only a tiny fraction of the total computational time. And, last but not least, the IPM and treecodes are not incompatible. On the contrary, they are complementary. Indeed, the IPM is perfectly suitable for combining with a hierarchical treecode and, therefore, for combining the gains provided by the two methods. IPM allows us to shoot far fewer rays, while hierarchical treecodes make the calculation for each ray much faster.

### 8.7.3 Summary

We have seen how the IRS can be improved in two different and complementary ways to speed up computations and make the calculation of magnification maps accessible for a standard desktop computer. A few final words may help the readers to decide which path to take if they need a magnification map. I try here to give a (far from precise) recipe:

- (i) For magnification maps of a few lenses IPM is, without doubt, the way to go. This is particularly true if we need to produce many of these maps in a reasonably short amount of time. If only a handful of maps are needed, maybe plain IRS is enough.
- (ii) If we need microlensing magnification maps that include many lenses, then a hierarchical treecode is the best choice. How many lenses are ‘many’ in this context? Well, there is no fixed number, but a few thousand lenses is a reasonable estimate



for when treecodes start to pay off. Close to this number, probably the IPM is as competitive as a treecode; well below that, IPM is certainly faster.

- (iii) If you need mass production of maps, maybe you should seriously consider getting a GPU and giving it a try.

Hopefully, very soon an IPM algorithm with a hierarchical treecode will be available and we will not need to make a choice between them.

## Acknowledgements

I wish to express my gratitude to the organizers of this Winter School, Evencio Mediavilla and José A. Muñoz, for organizing such a wonderful school and for inviting me to deliver the lectures. In fact, I am doubly indebted to Evencio, as he has taught me most of what I know about lensing in general and making microlensing maps in particular.

I also want to thank Simon Verley for introducing me to **Python** and for having the patience to answer my endless questions about it during my first steps with it.

## REFERENCES

- Alcock, C. et al. 2000, *ApJ*, **541**, 270
- Barnes, J. & Hut, P. 1986, *Nature*, **324**, 446
- Bate, N. F., Fluke, C. J., Barsdell, B. R., Garsden, H. & Lewis, G. F. 2010, *New Astron.*, **15**, 726
- Chang, K. & Refsdal, S. 1979, *Nature*, **282**, 561
- Chang, K. & Refsdal, S. 1984, *A&A*, **132**, 168
- Jiménez-Vicente, J., Mediavilla, E., Muñoz, J. A. & Kochanek, C. S. 2012, *ApJ*, **751**, 106
- Katz, N., Balbus, S. & Paczynski, B. 1986, *A&A*, **306**, 2
- Kayser, R., Refsdal, S. & Stabell, R. 1986, *A&A*, **166**, 36
- Kochanek, C. S. 2004, *ApJ*, **605**, 58
- Mediavilla, E., Muñoz, J. A., López, P., Mediavilla, T., Abajas, C., Gonzalez-Morcillo, C. & Gil-Merino, R. 2006, *ApJ*, **653**, 942
- Mediavilla, E., Muñoz, J. A., Falco, E., Motta, V., Guerras, E., Canovas, H., Jean, C. Osoz, A. & Mosquera, A. 2009, *ApJ*, **706**, 1451
- Mediavilla, E., Mediavilla, T., Muñoz, J. A., Ariza, O., López, P., González-Morcillo, C. & Jiménez-Vicente, J. 2011, *ApJ*, **741**, 42
- Mortonson, M. J., Schechter, P. L. & Wambsganss, J. 2005, *ApJ*, **628**, 594
- Muñoz, J. A., Mediavilla, E., Kochanek, C. S., Falco, E. E. & Mosquera, A. M. 2011, *ApJ*, **742**, 67
- Schneider, P., Ehlers, J. & Falco, E. E. 1999, Astronomy and Astrophysics Library, *Gravitational Lenses* (Berlin: Springer-Verlag)
- Schneider, P., Kochanek, C. & Wambsganss, J. 2006, Saas-Fee Advanced Course 33, *Gravitational Lensing: Strong, Weak and Micro* (Berlin: Springer-Verlag)
- Schneider, P. & Weiss, A. 1986, *A&A*, **164**, 237
- Schneider, P. & Weiss, A. 1987, *A&A*, **171**, 49
- Shakura, N. I. & Sunyaev, R. A. 1973, *A&A*, **24**, 337
- Shipman, J. W. 2011, *A Python programming tutorial*, <http://infohost.nmt.edu/tcc/help/pubs/lang/pytut1/>
- Wambsganss, J. 1990, PhD Thesis, Munich, available as MPA Report 550
- Wambsganss, J. 1999, *J. Comp. App. Math.*, **199**, 353